Thinking about high dimensional data

Machine Learning for Single Cell Analysis

Day 1 – Morning session

What does “dimensionality” mean?

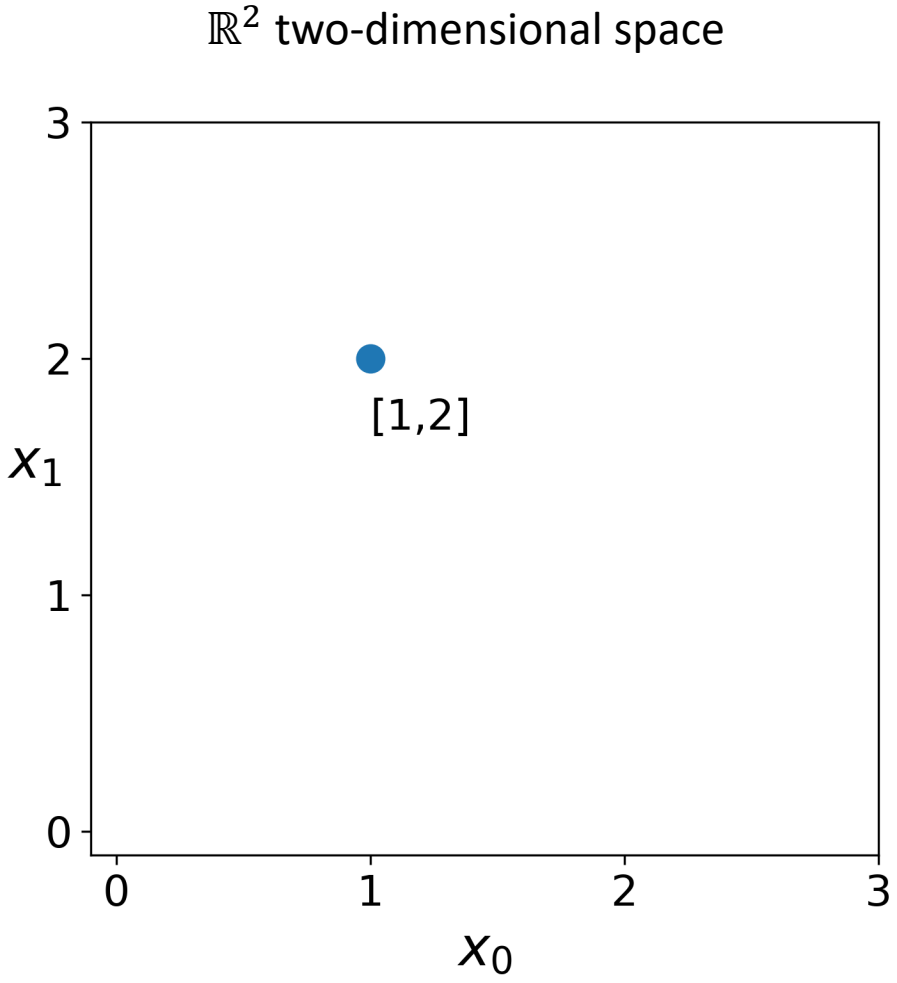

What does “dimensionality” mean?

vector coordinate space

$\boldsymbol{a}_0 \in \mathbb{R}^2$

$\boldsymbol{a}_0 = [1,2]$

x_0 x_1

features

$\boldsymbol{a}_{[:,0]}$

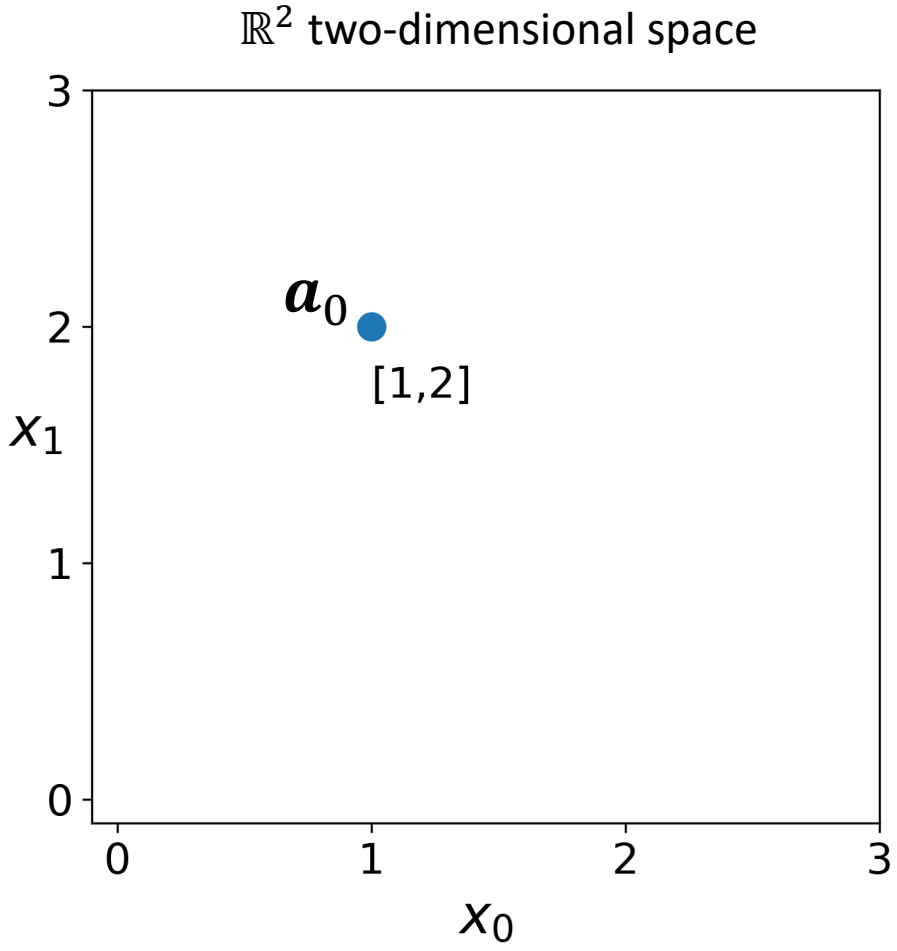

What does “dimensionality” mean?

vector coordinate space

$\boldsymbol{a}_0 \in \mathbb{R}^2$

$\boldsymbol{a}_0 = [1,2]$

x_0 x_1

features

$\mathbb{R}^2$ two-dimensional space

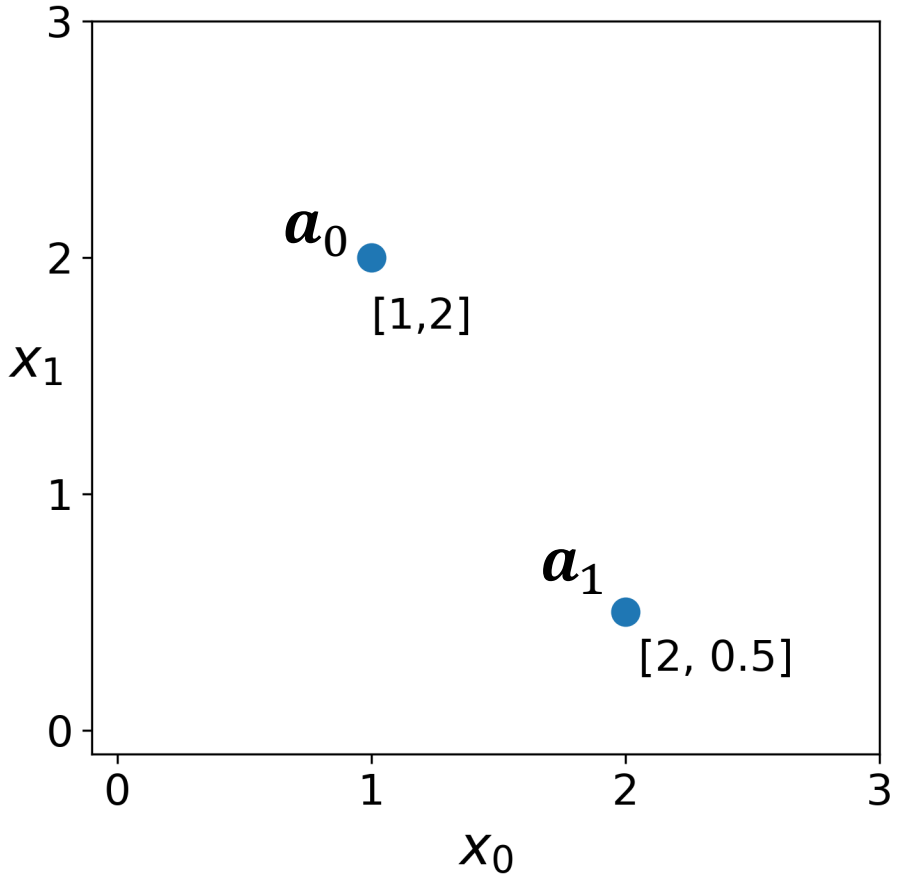

$\boldsymbol{a}_{[:,0]}$

What does “dimensionality” mean?

$\boldsymbol{a} \in \mathbb{R}^2$

$\boldsymbol{a} = [1,2]$

x_0 x_1

$\boldsymbol{b} \in \mathbb{R}^3$

$\boldsymbol{b} = [2,4,1]$

x_0 x_1 x_2

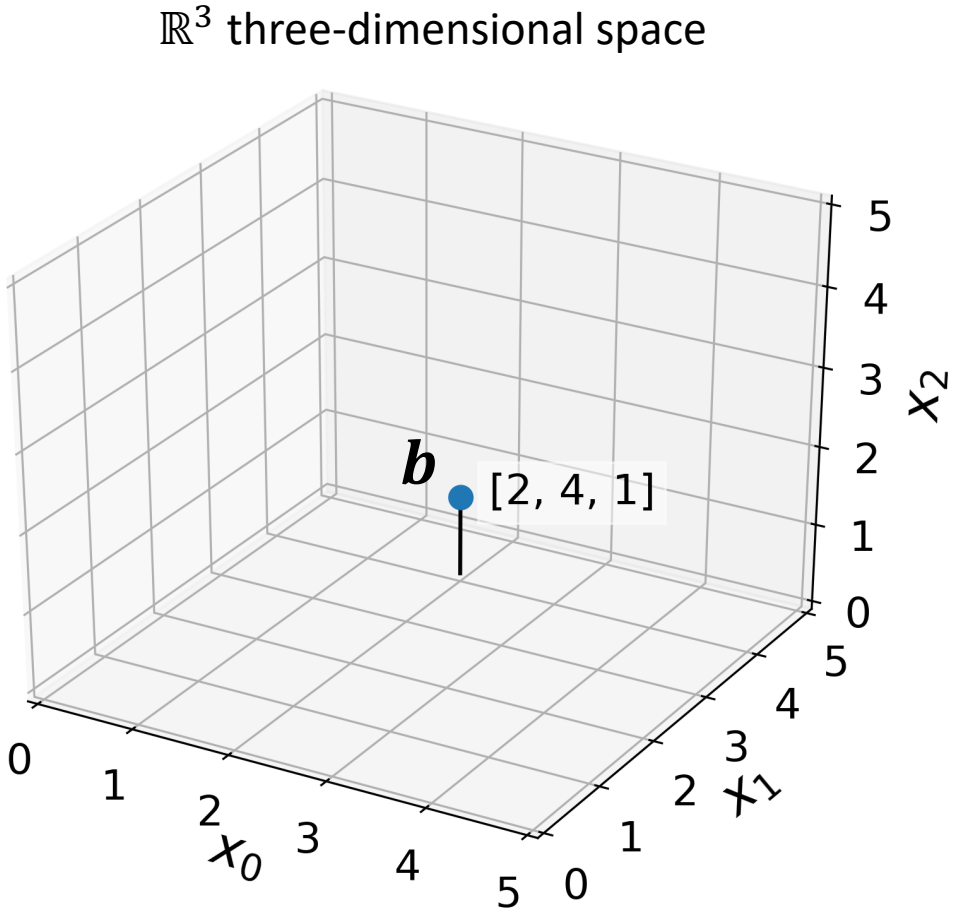

What does “dimensionality” mean?

$\boldsymbol{a} \in \mathbb{R}^2$

$\boldsymbol{a} = [1,2]$

x_0 x_1

$\boldsymbol{b} \in \mathbb{R}^3$

$\boldsymbol{b} = [2,4,1]$

x_0 x_1 x_2

$\boldsymbol{c} \in \mathbb{R}^3$

$\boldsymbol{c} = [2,4,1,1]$

x_0 x_1 x_2 x_3

$\mathbb{R}^4$ four-dimensional space

???

What does “dimensionality” mean?

$\boldsymbol{a} \in \mathbb{R}^2$

$\boldsymbol{a} = [1,2]$

x_0 x_1

$\boldsymbol{b} \in \mathbb{R}^3$

$\boldsymbol{b} = [2,4,1]$

x_0 x_1 x_2

$\boldsymbol{c} \in \mathbb{R}^4$

$\boldsymbol{c} = [2,4,1,1]$

x_0 x_1 x_2 x_3

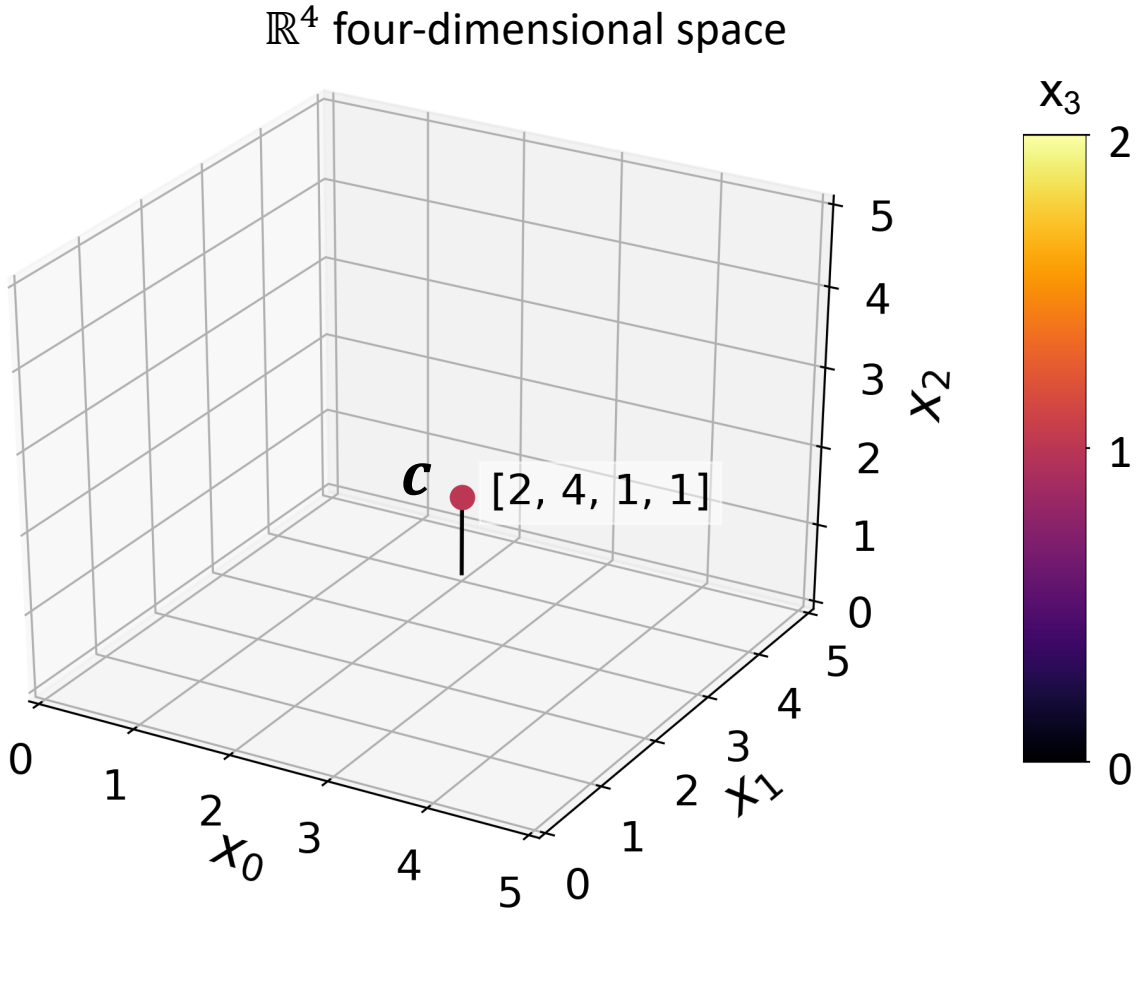

What does “dimensionality” mean?

$\boldsymbol{a} \in \mathbb{R}^2$

$\boldsymbol{a} = [1,2]$

x_1 x_2

$\boldsymbol{b} \in \mathbb{R}^3$

$\boldsymbol{b} = [2,4,1]$

x_1 x_2 x_3

$\boldsymbol{c} \in \mathbb{R}^4$

$\boldsymbol{c} = [2,4,1,1]$

x_1 x_2 x_3 x_4

$\mathbb{R}^5$ five-dimensional space

???

What do x_1 x_2, …, x_n represent?

We describe coordinate spaces of n dimensions as $\mathbb{R}^n$

$\boldsymbol{a} =$

	alcohol	hue
Wine0	14.23	1.04

$\boldsymbol{a}_0$ $\boldsymbol{a}_1$

$\boldsymbol{b} =$

	Age	Height (m)	Weight (kg)
Player0	23	2.31	91

$\boldsymbol{b}_0$ $\boldsymbol{b}_1$ $\boldsymbol{b}_2$

$\boldsymbol{c} =$

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
Iris0	5.1	3.5	1.4	0.2

UCI ML Wine Data Set

What if we have more than one observation?

$\boldsymbol{a} =$

	alcohol	hue
Wine0	14.23	1.04

$\boldsymbol{a}_0$ $\boldsymbol{a}_1$

columns (features)

$\boldsymbol{A} =$ matrix

	alcohol	hue
Wine0	14.23	1.04
Wine1	13.20	1.05
Wine2	13.16	1.03
Wine3	14.37	0.86
Wine4	13.24	1.04
Wine5	14.20	1.05
Wine6	14.39	1.02
Wine7	14.06	1.06
Wine8	14.83	1.08
Wine9	13.86	1.01

rows (observations)

UCI ML Wine Data Set

What if we have more than one observation?

$\boldsymbol{a} =$

	alcohol	hue
Wine0	14.23	1.04

$\boldsymbol{a}_0$ $\boldsymbol{a}_1$

columns (features)

$\boldsymbol{A} =$ matrix

	alcohol	hue
Wine0	14.23	1.04
Wine1	13.20	1.05
Wine2	13.16	1.03
Wine3	14.37	0.86
Wine4	13.24	1.04
Wine5	14.20	1.05
Wine6	14.39	1.02
Wine7	14.06	1.06
Wine8	14.83	1.08
Wine9	13.86	1.01

$\epsilon\ \mathbb{R}^2$

UCI ML Wine Data Set

How do we reference specific observations in a matrix?

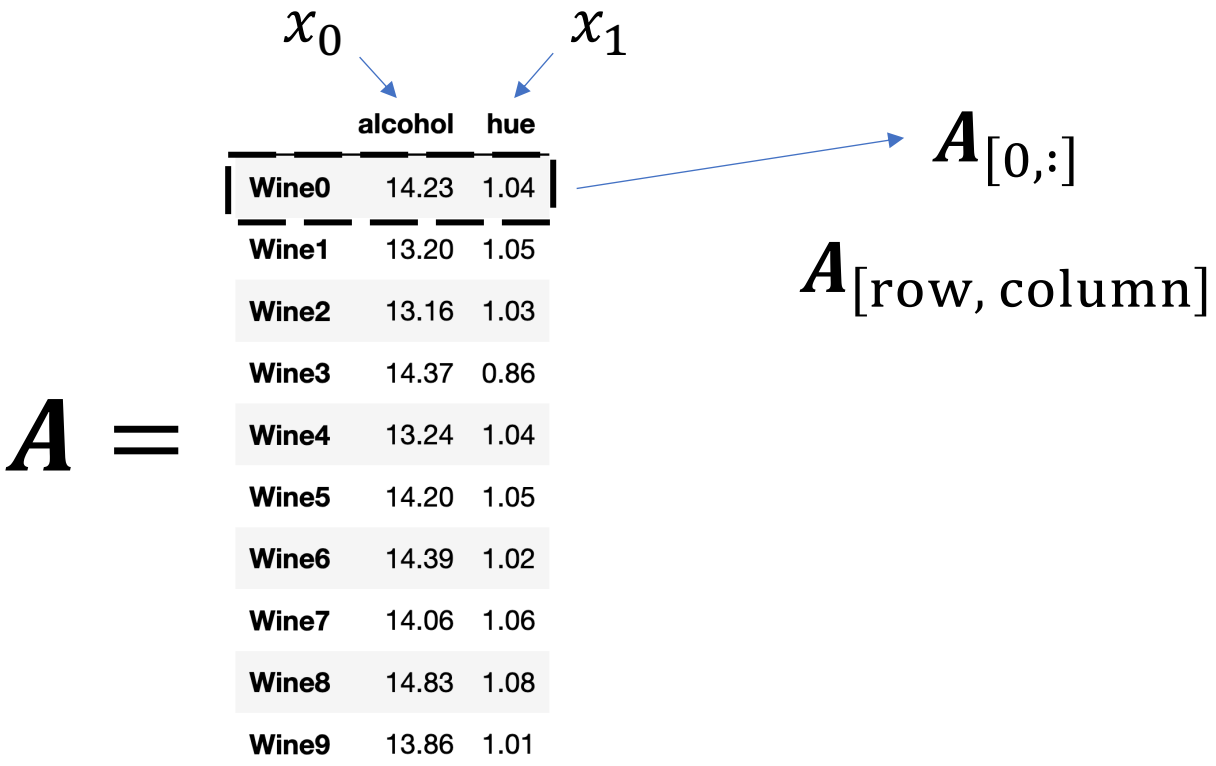

What if we have more than one observation?

We can describe coordinate spaces of n dimensions as $\mathbb{R}^n$

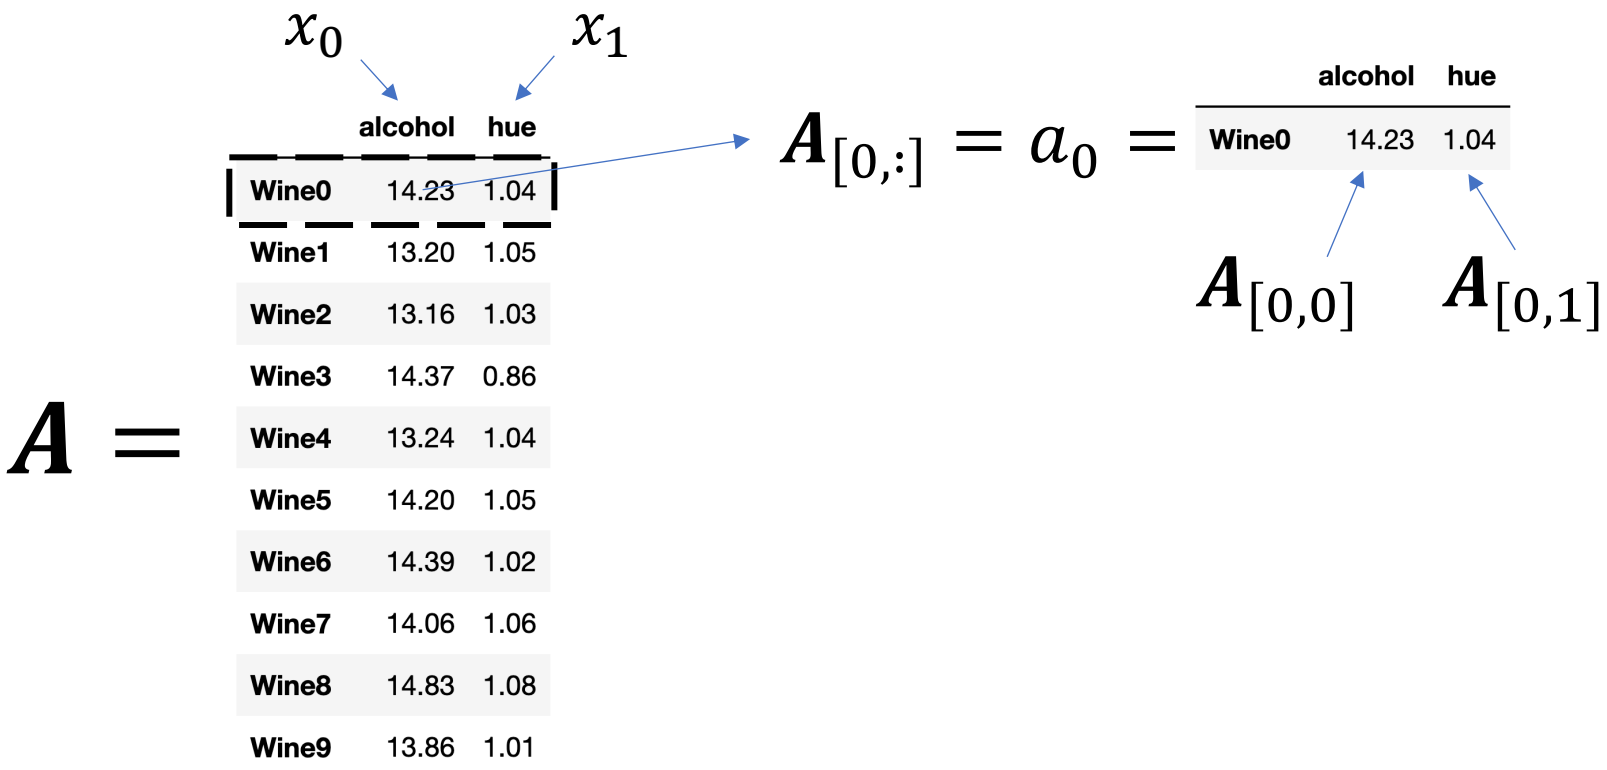

What if we have more than one observation?

We can describe coordinate spaces of n dimensions as $\mathbb{R}^n$

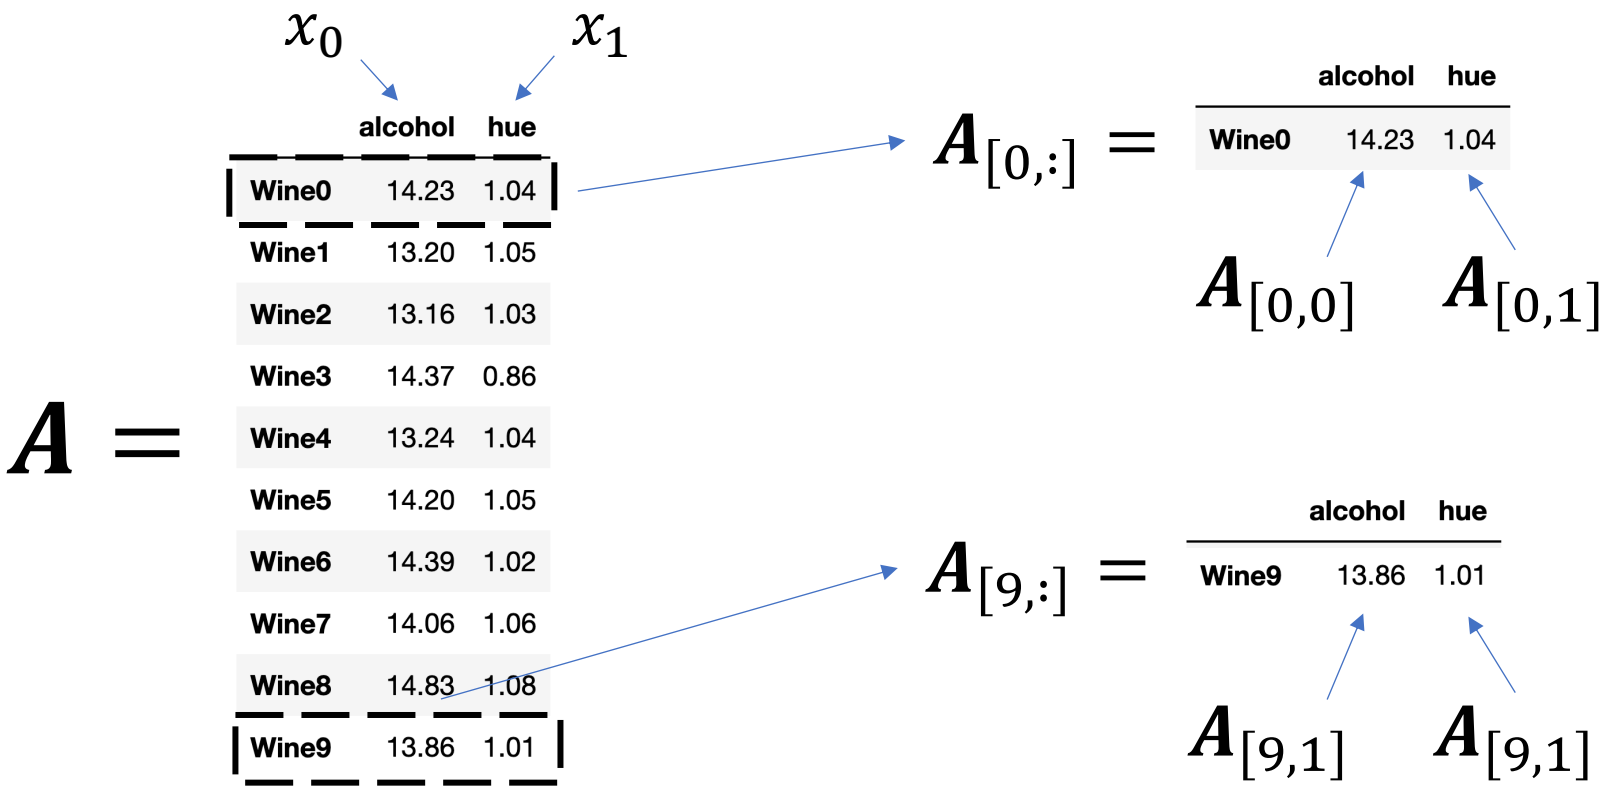

Exercise!

Loading and plotting the UCI wine dataset

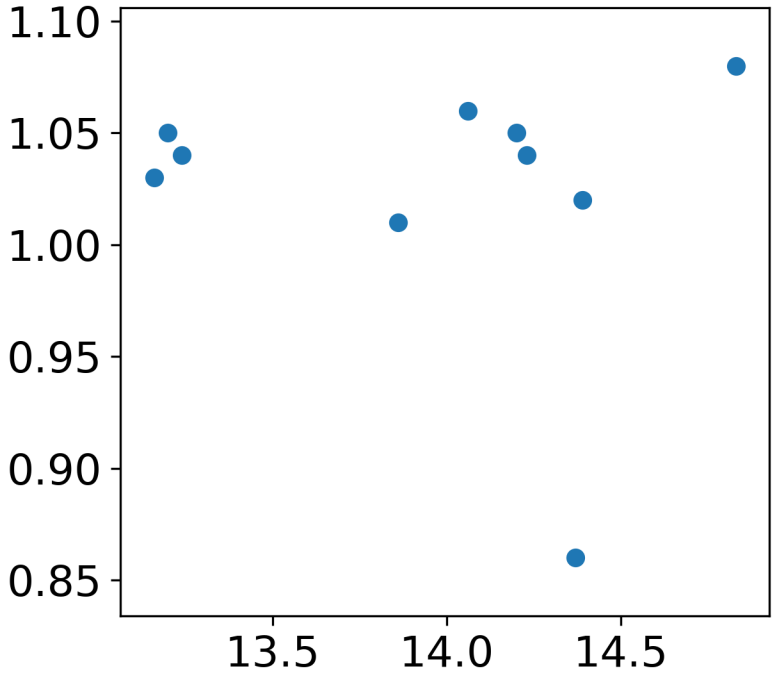

Exercise!

Loading and plotting the UCI wine dataset

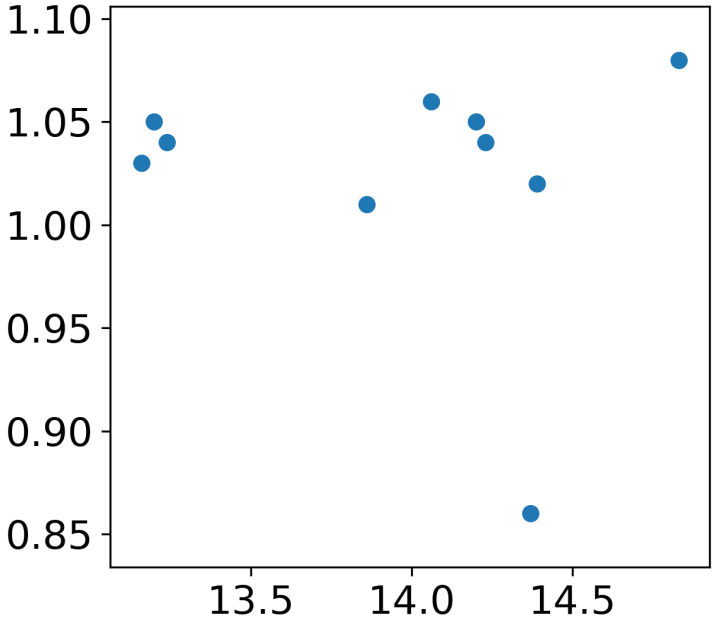

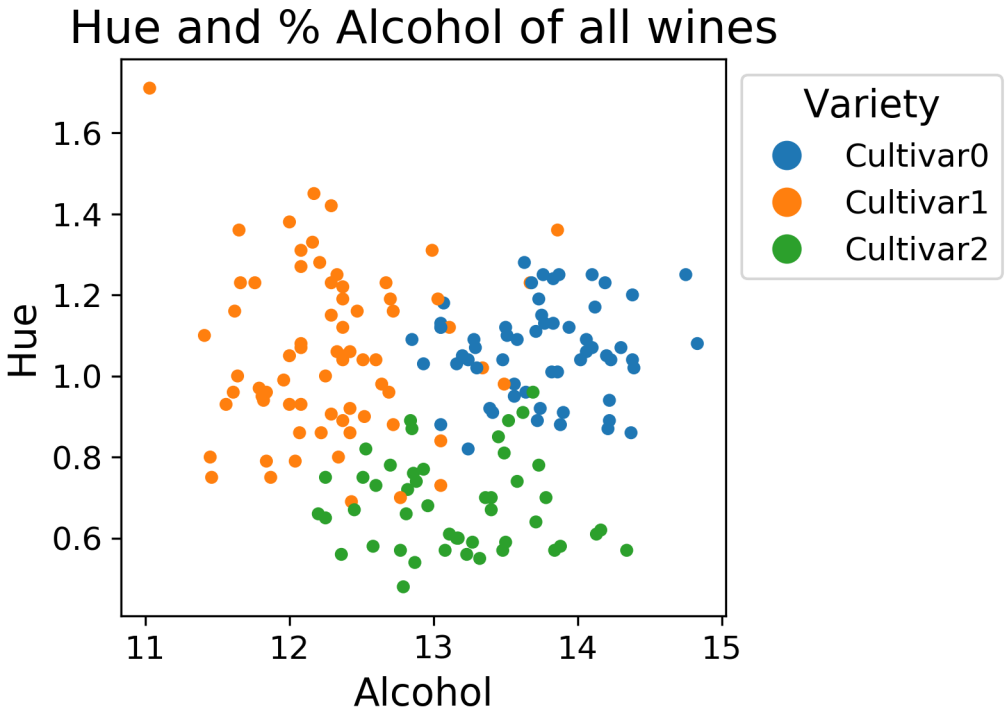

Introducing scprep

- Code: https://github.com/KrishnaswamyLab/scprep
- Docs: https://scprep.readthedocs.io

- Modules
  - filter
  - io
  - measure
  - normalize
  - **plot**
  - reduce
  - run
  - select
  - stats
  - transform
  - utils

Plotting functions:

- histogram
- jitter
- marker
- marker_plot
- plot_gene_set_expression
- plot_library_size
- rotate_scatter3d
- scatter
- **scatter2d** = `scprep.plot.scatter2d`
- scatter3d
- scree
- scree_plot

Exercise!

Loading and plotting the UCI wine dataset

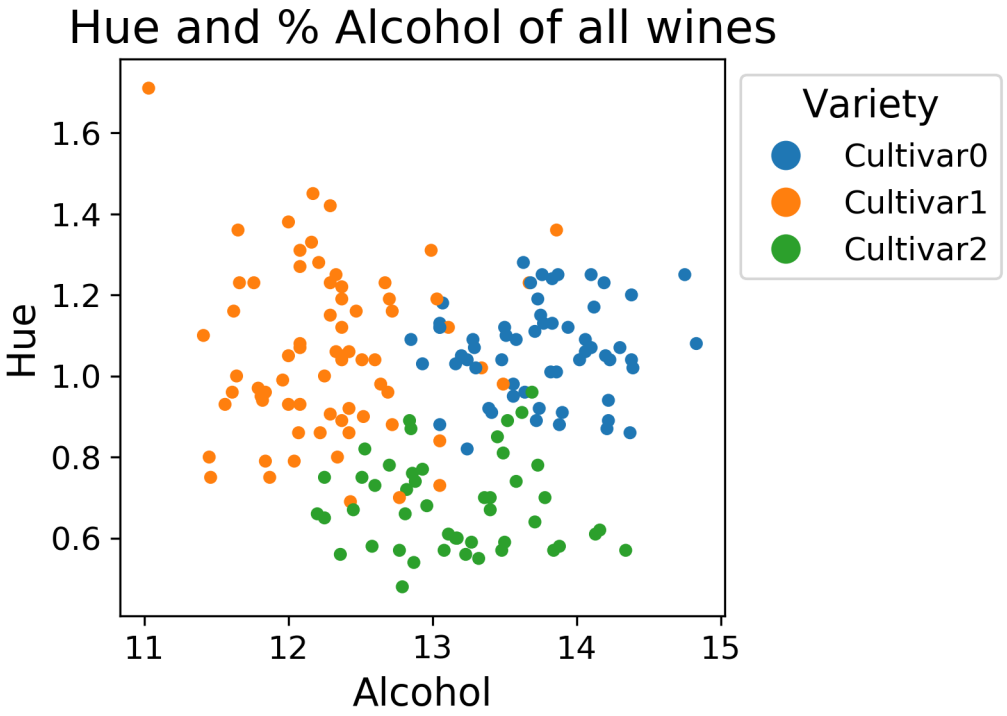

How do features differentiate groups within a dataset?

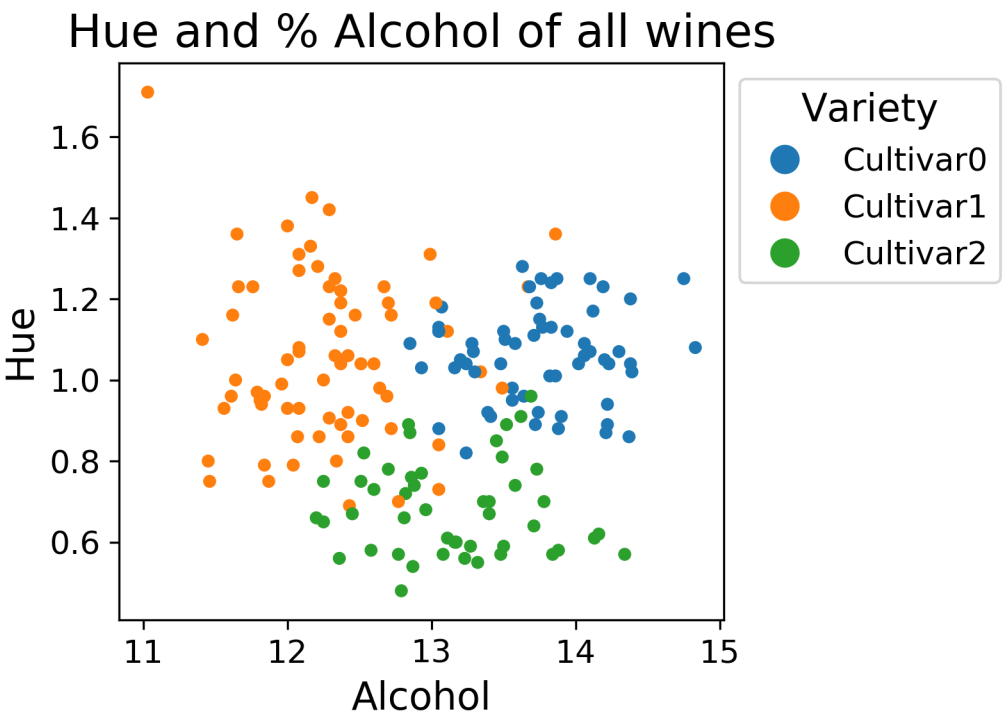

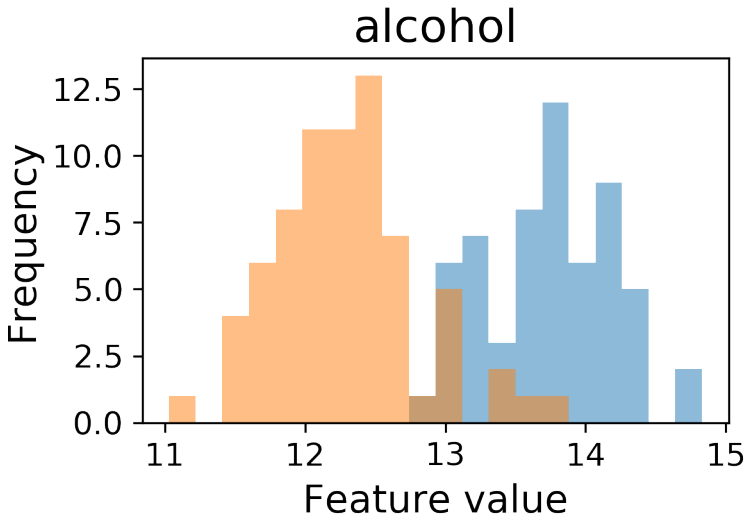

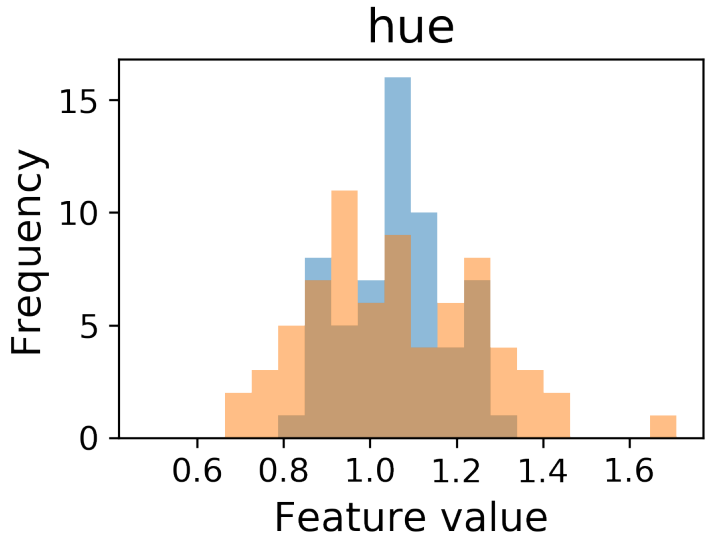

Introduction to Manifold Learning

How far away is each point from $\boldsymbol{a}_0$?

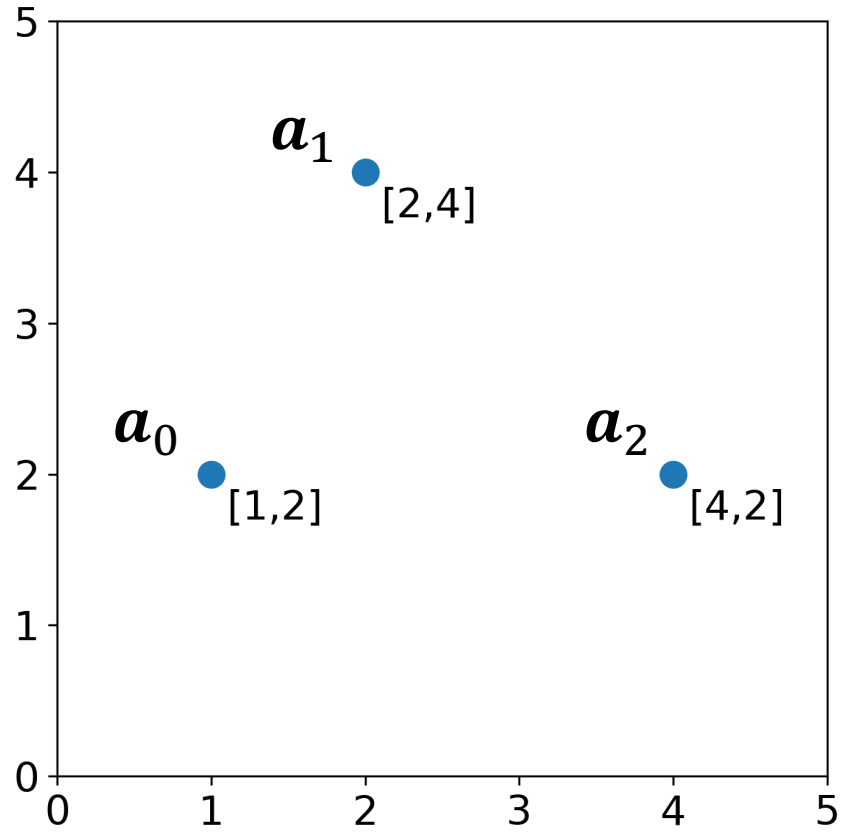

How far away is each point from a_0?

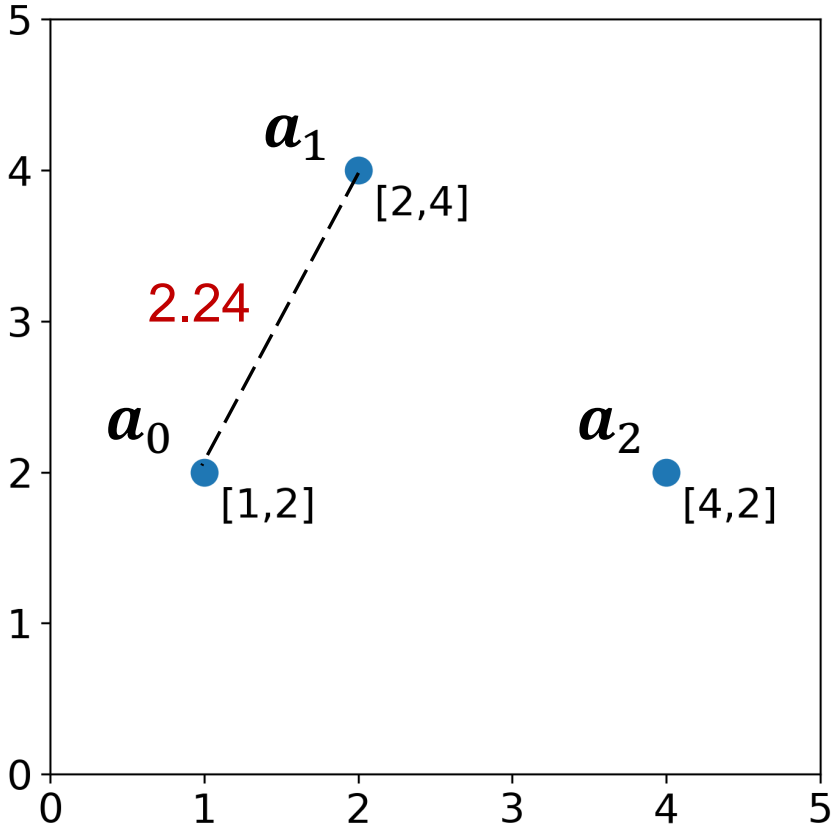

Euclidean distance

$$d_{euclidean}(a_0, a_1) = \|a_0 - a_1\|_2^2$$

$$= \sqrt{(a_{0,0} - a_{1,0})^2 + (a_{0,1} - a_{1,1})^2}$$

$$= \sqrt{(1 - 2)^2 + (2 - 4)^2}$$

$$\approx 2.24$$

How far away is each point from a_0?

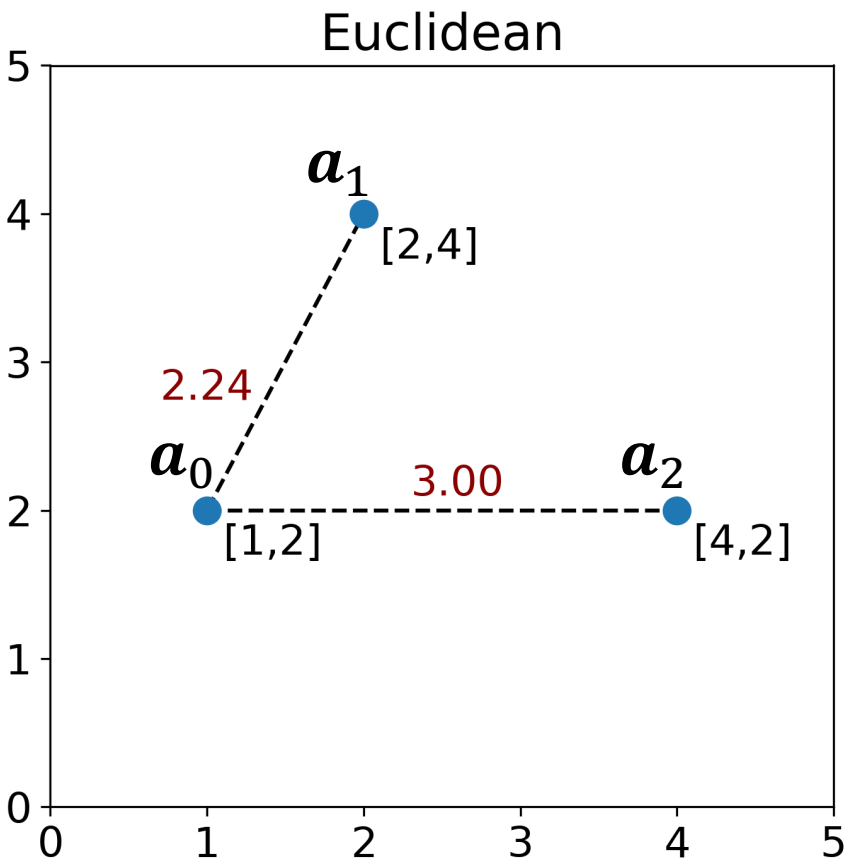

How far away is each point from $\boldsymbol{a}_0$?

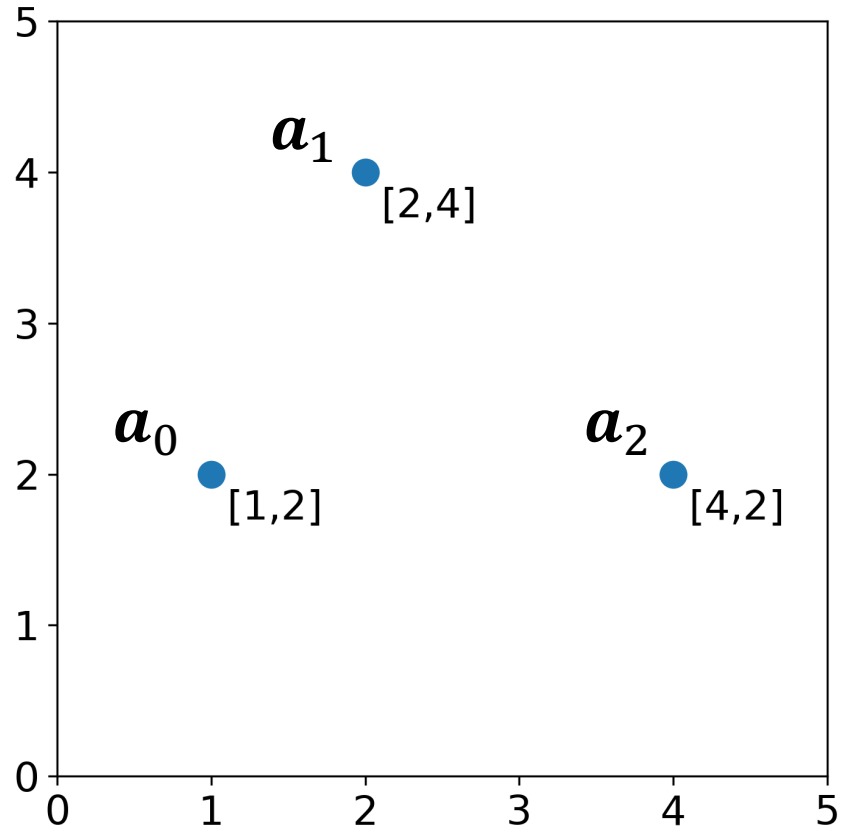

How far away is each point from a_0?

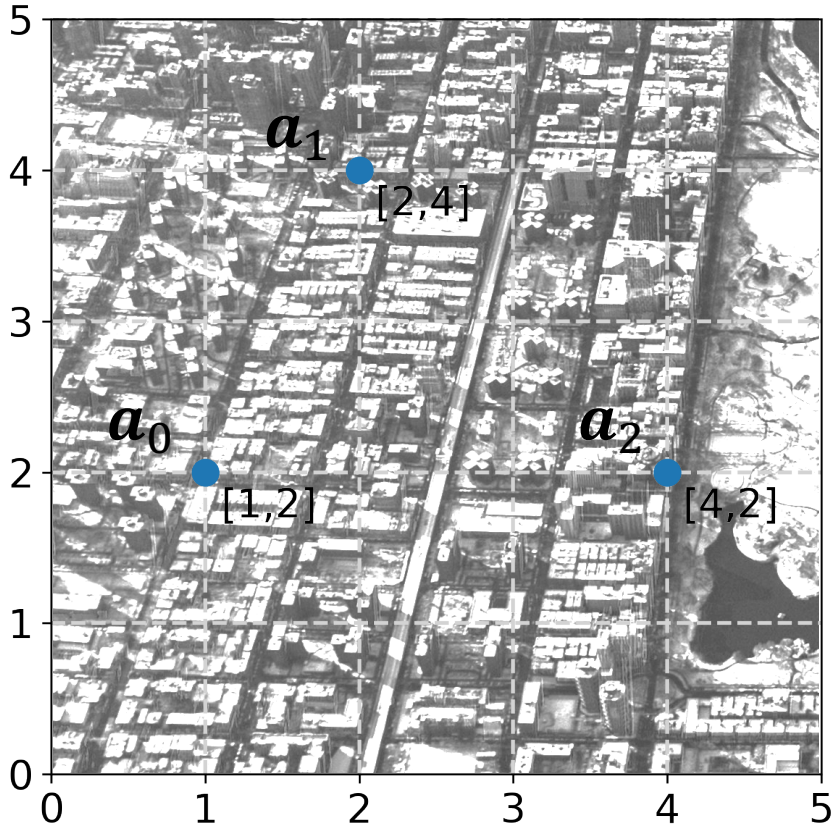

How far away is each point from $\boldsymbol{a}_0$?

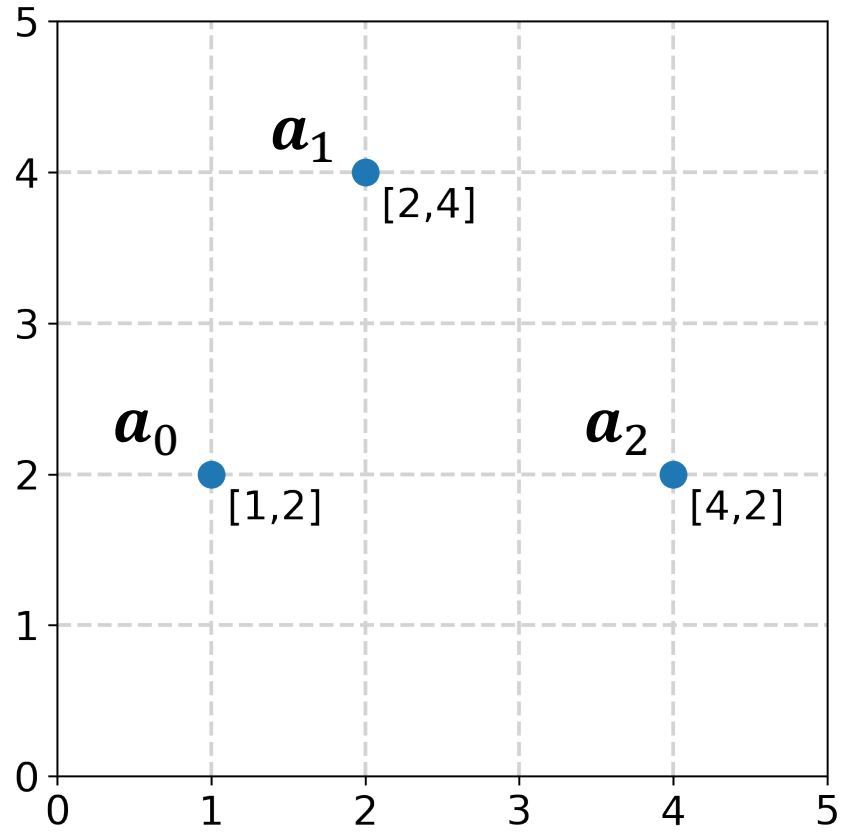

How far away is each point from a_0?

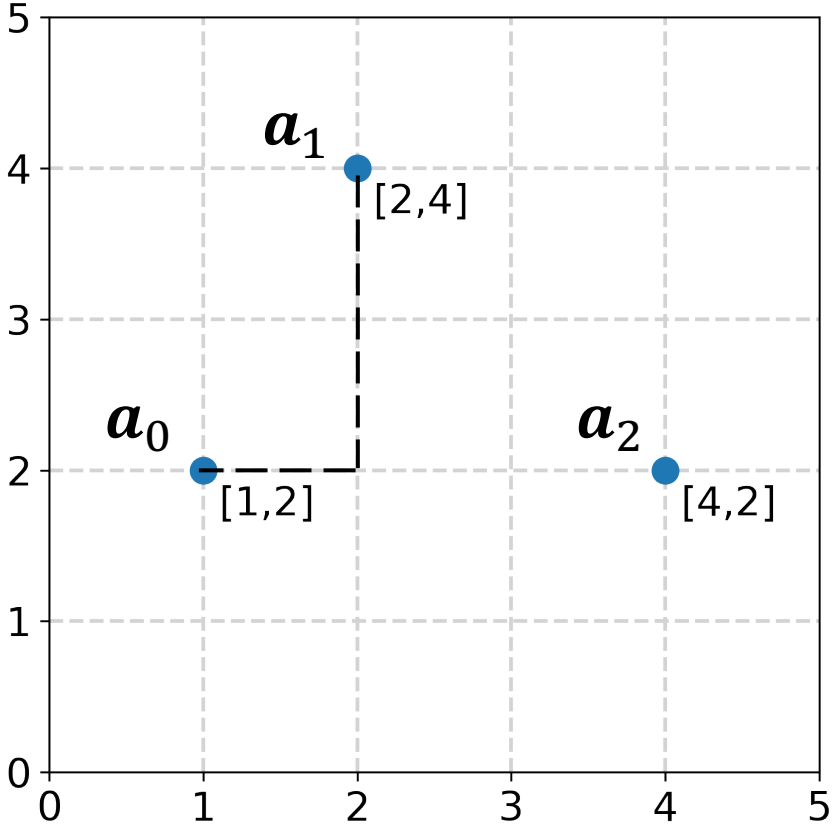

Manhattan distance

$$d_{manhattan}(a_0, a_1) = |a_{0,1} - a_{1,1}|$$

$$= |a_{0,0} - a_{1,0}| + |a_{0,1} - a_{1,1}|$$

$$= |1 - 2| + |2 - 4|$$

$$= 3$$

How far away is each point from a_0?

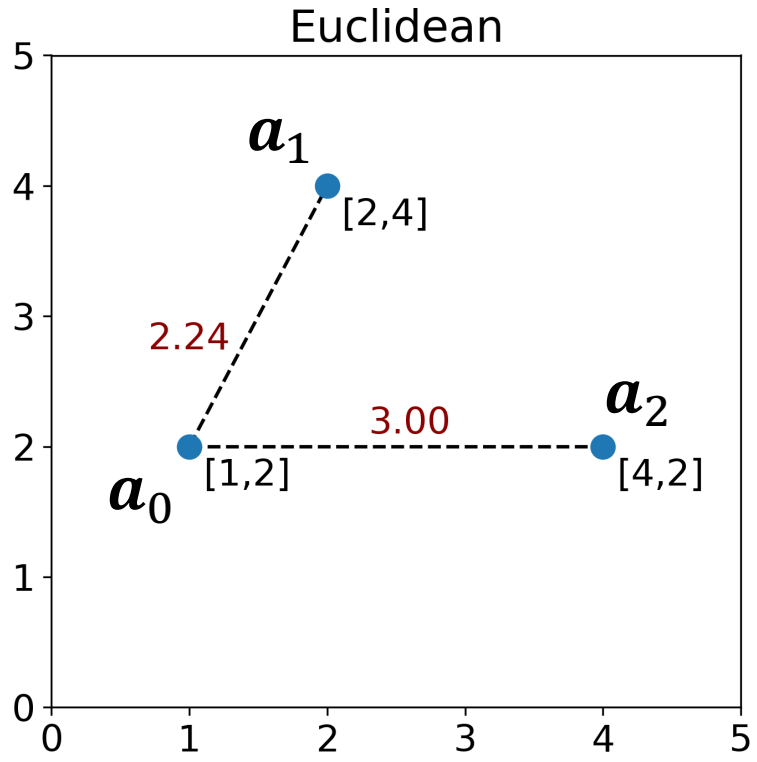

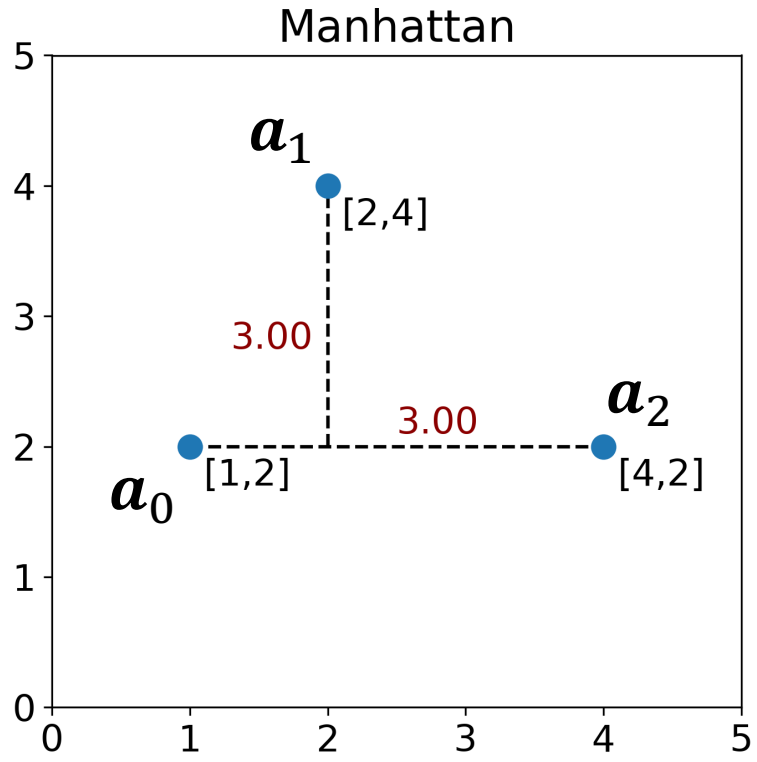

Exercise!

Calculating distances

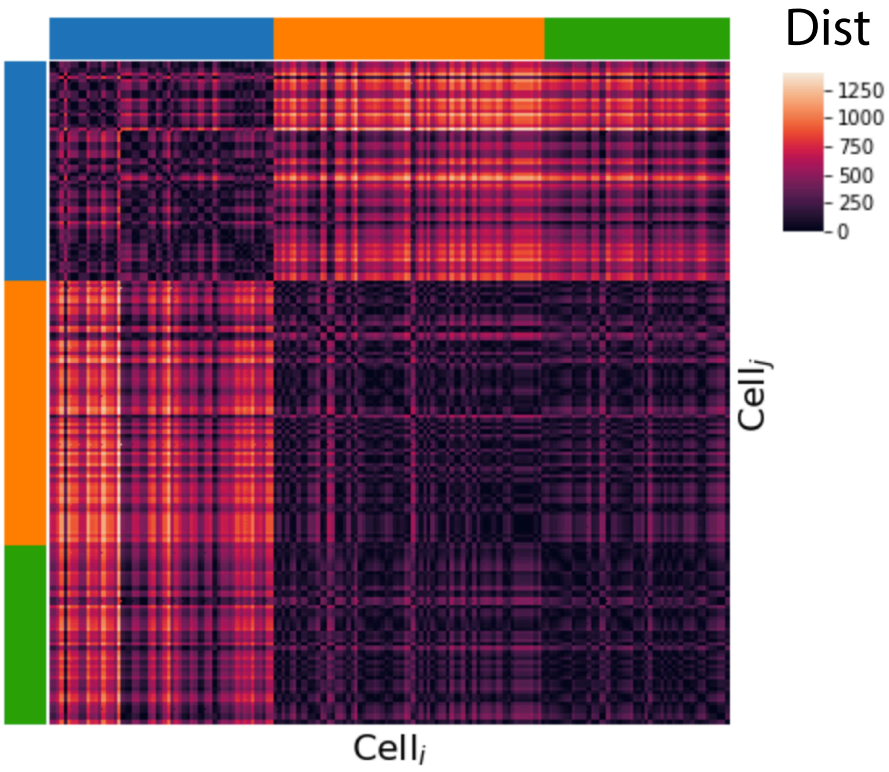

Exercise!

Calculating distances

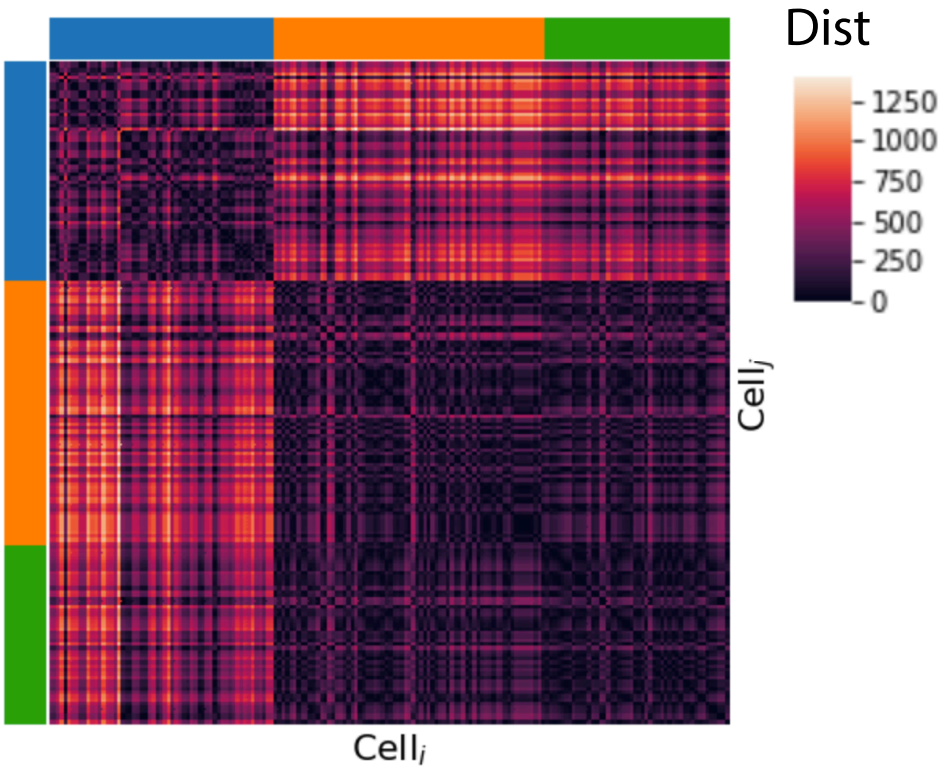

Distance defines **structure** in a dataset

- How **different** are two data points?
- How many **types** exist in my dataset?
- Which types are **closest** to each other?
- Which **features** vary between dissimilar groups of data?

What is a manifold?

- Locally smooth
- Locally Euclidean
- Generally, lower dimensional than the ambient space (i.e. a subspace)

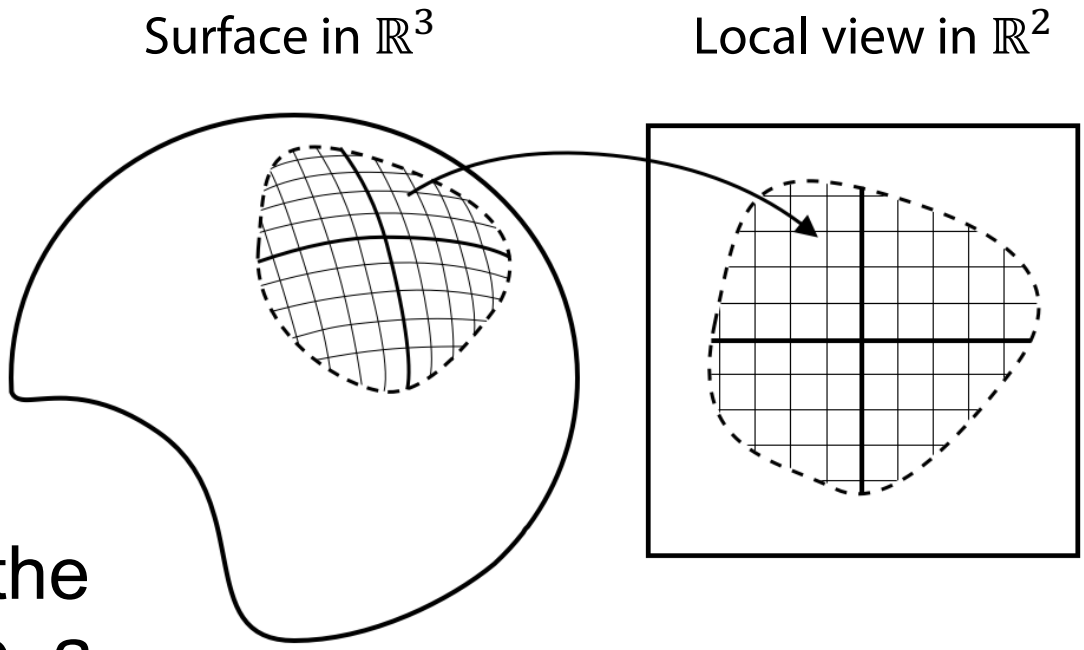

Examples of manifolds

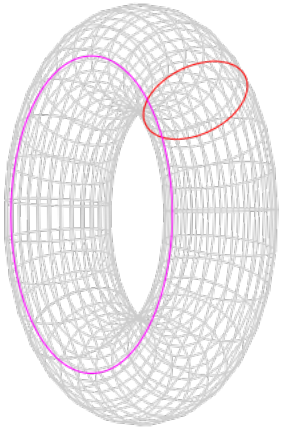

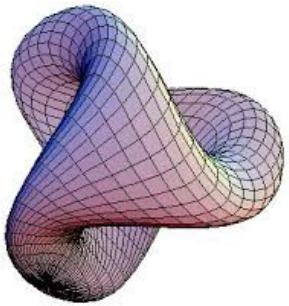

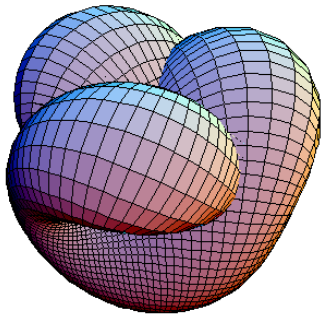

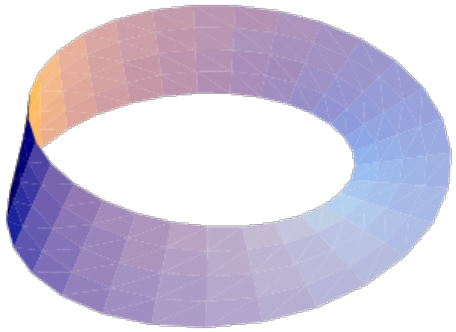

Examples of manifolds

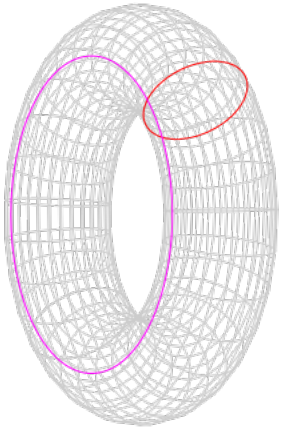

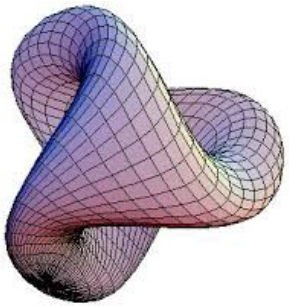

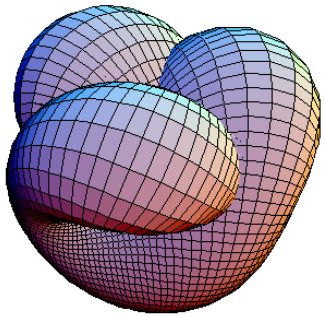

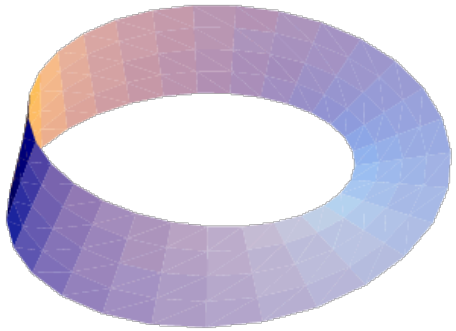

Waddington's landscape

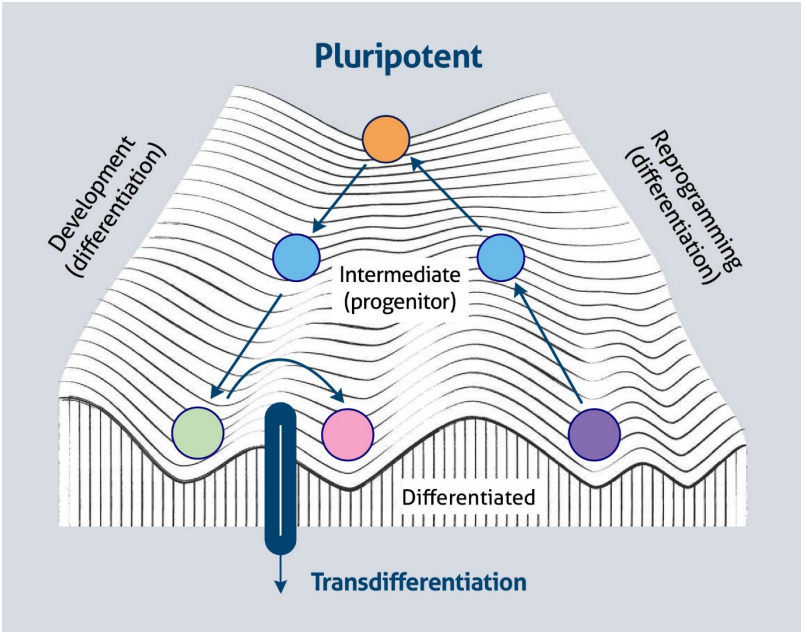

Examples of manifolds

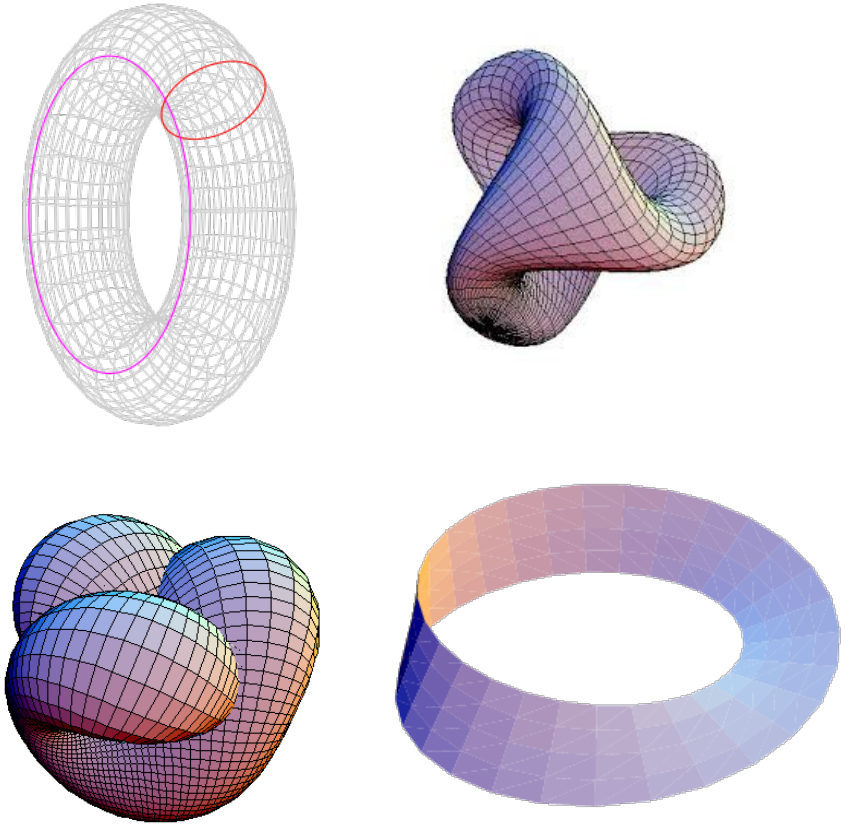

Waddington's landscape

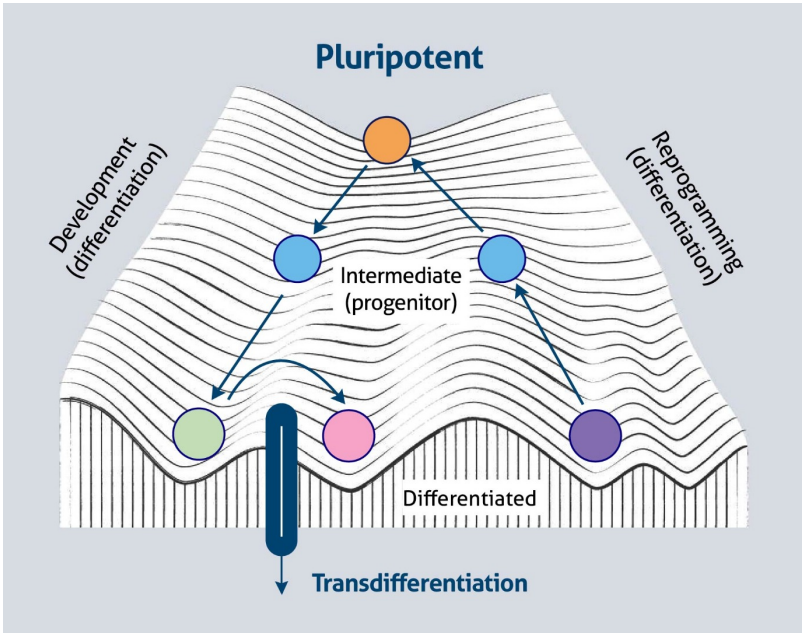

- Smooth transitions between states
- Small changes in gene expression are linear

NHV

⬇

MEX


East Siberian Sea
Arctic Ocean
Arctic Ocean
Bering Sea
Greenland Sea
Baffin Bay
Sweden
Norway
Belarus
Iceland
Northwestern Passages
Greenland
Poland
Canada
United Kingdom
Romania
Hudson Bay
Ireland
Turkey
Labrador Sea
France
Italy
Spain
Mediterranean Sea
Tunisia
Egypt
Portugal
Libya
United States
Algeria
North Atlantic Ocean
Chad
Western Sahara
Mali
Mauritania
Mexico
Gulf of Mexico
Nigeria
Cuba
Guinea
Ghana
Puerto Rico
Gabon
Guatemala
Caribbean Sea
Nicaragua
Venezuela
Guyana
Colombia
Suriname
Ecuador
Brazil
Peru
Bolivia
South Atlantic Ocean
Paraguay
Chile
Uruguay
Argentina

How far is NHV -> MEX?

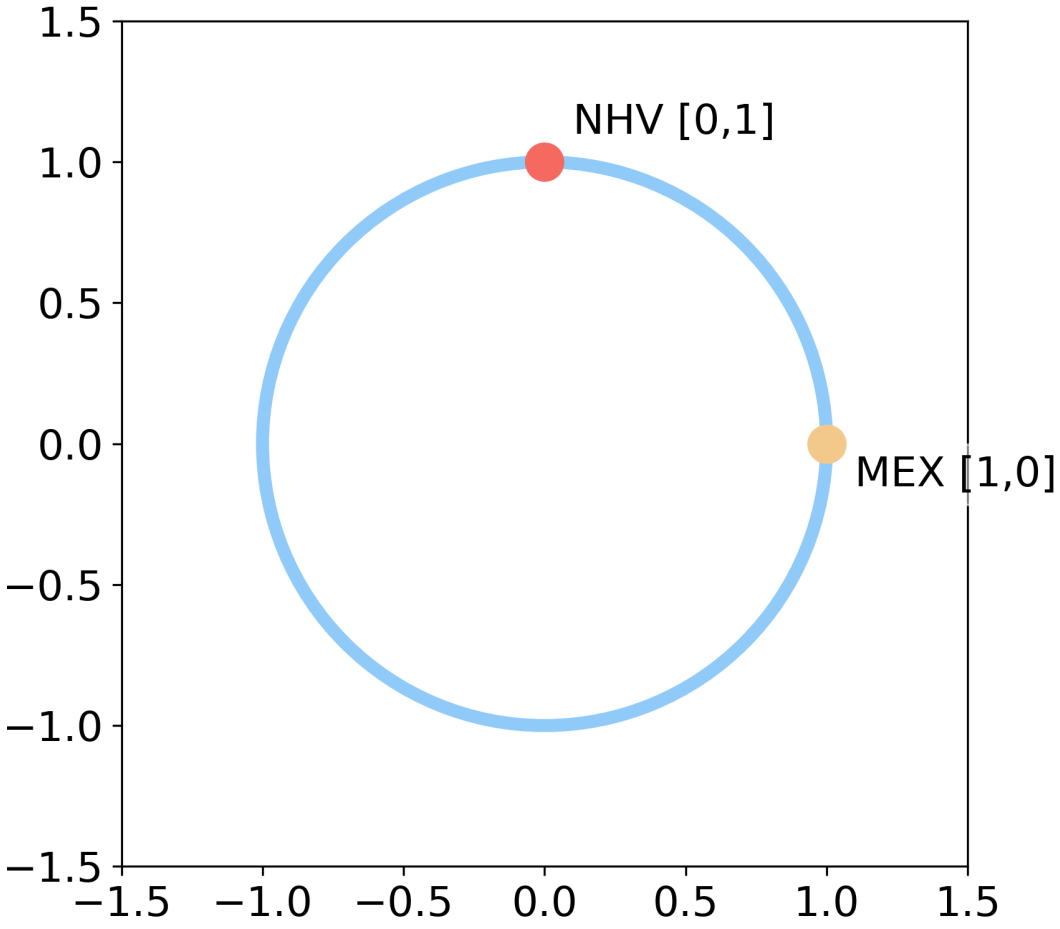

How far is NHV -> MEX?

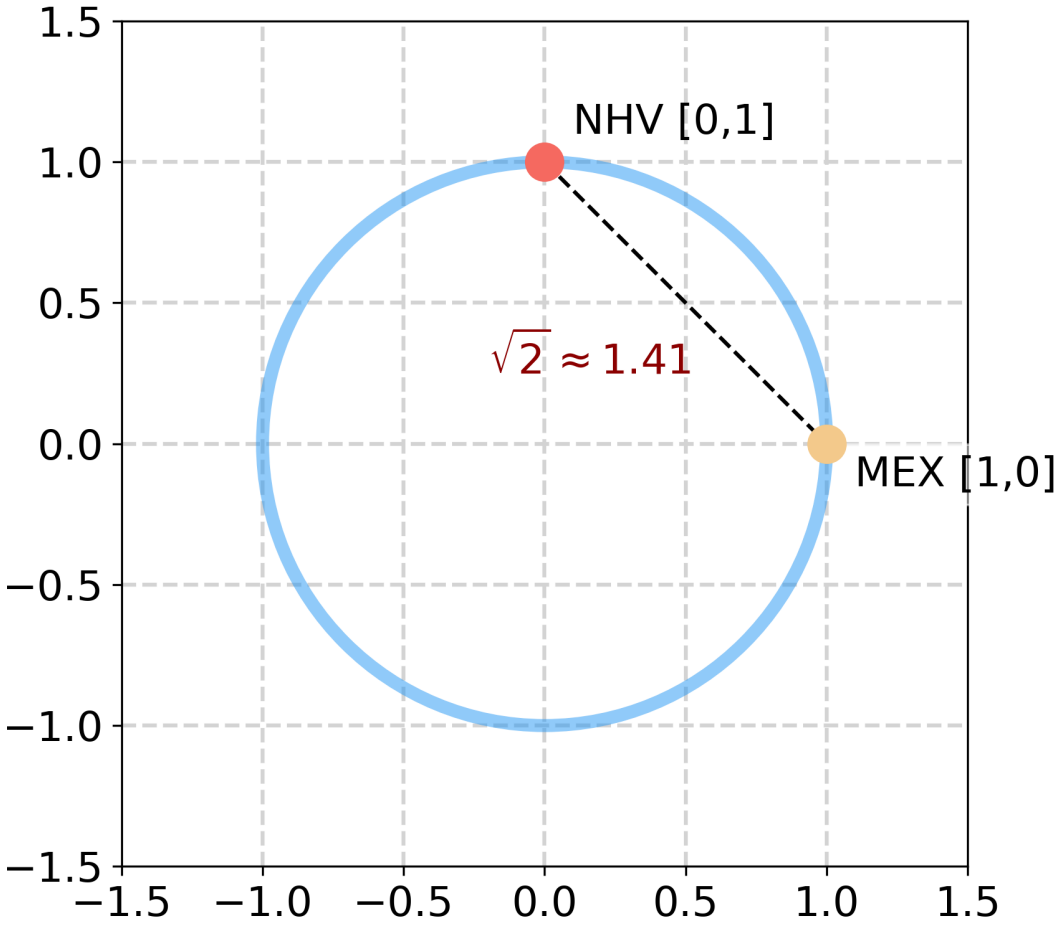

How far is NHV -> MEX?

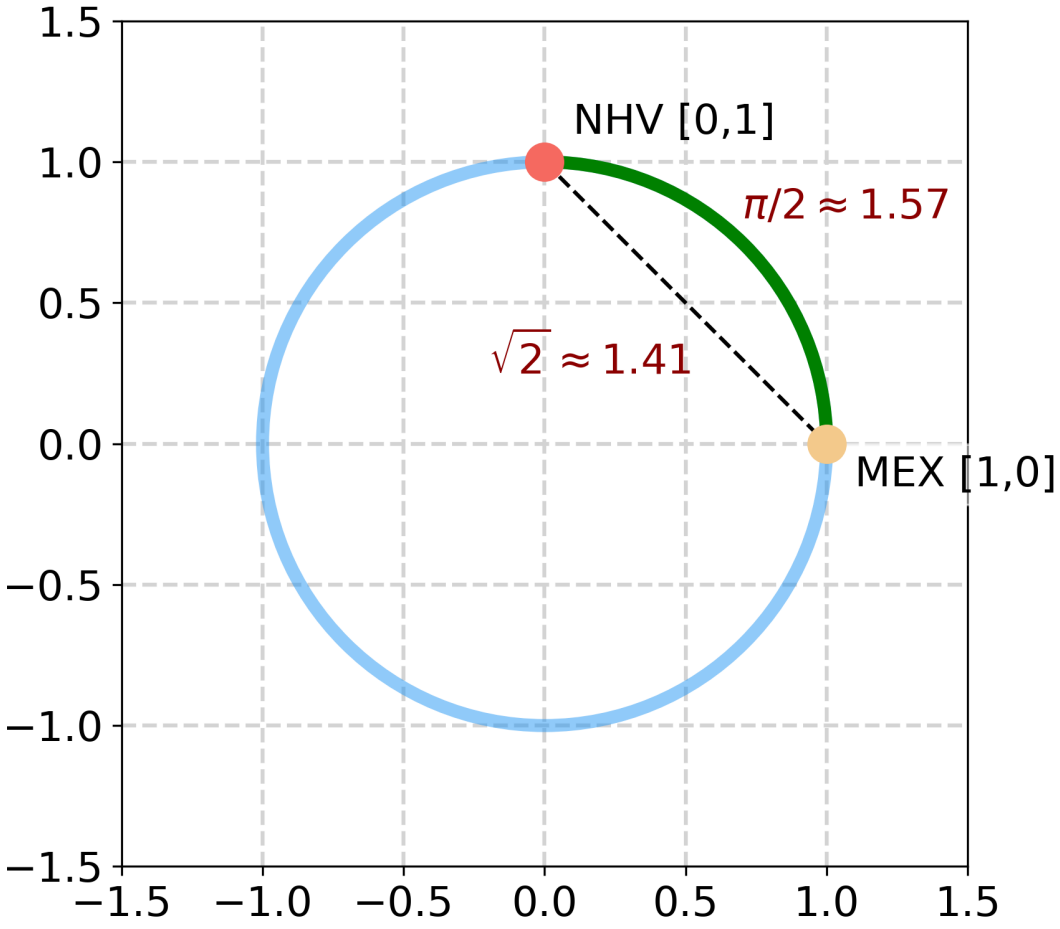

How far is NHV -> MEX?

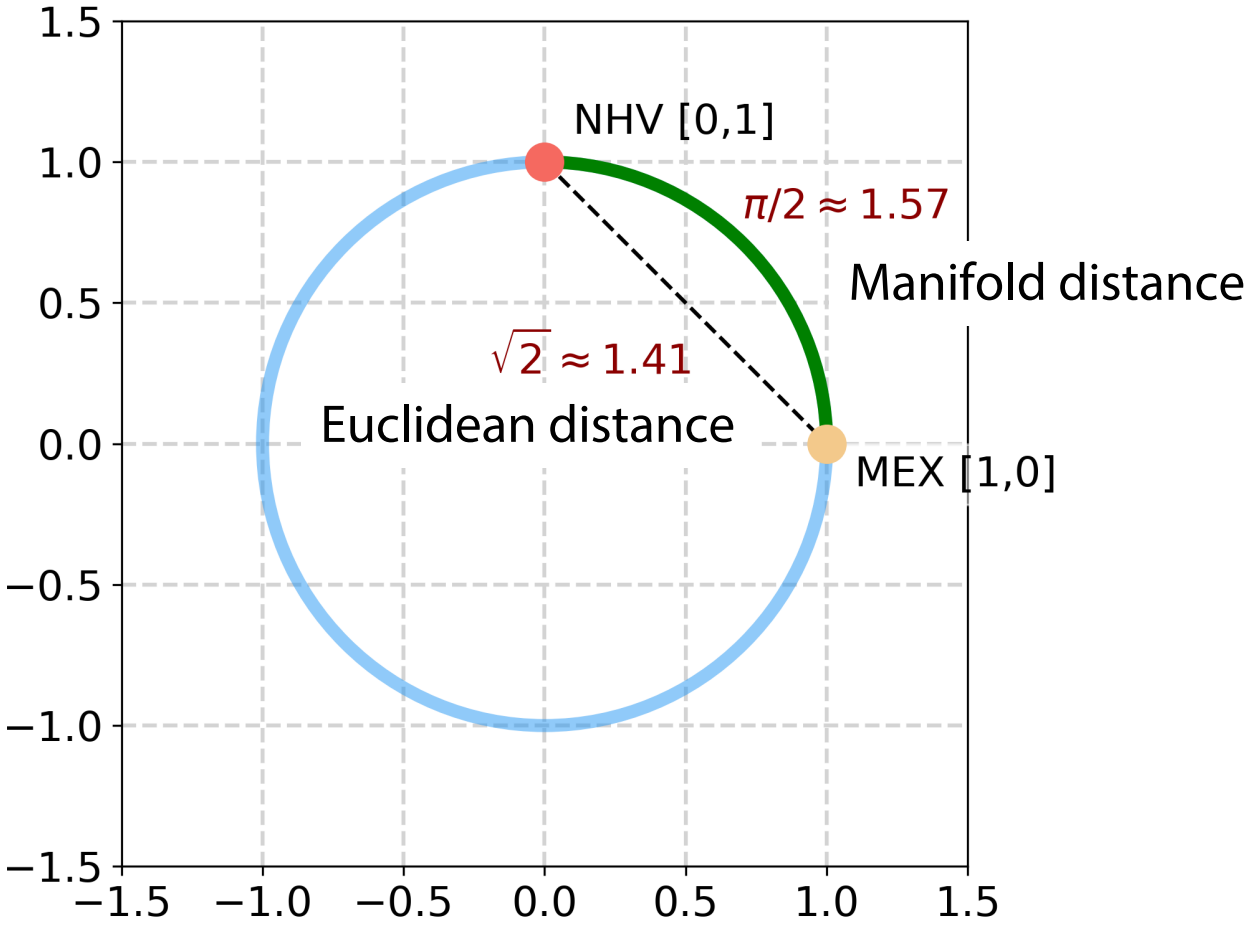

Is a circle locally Euclidean?

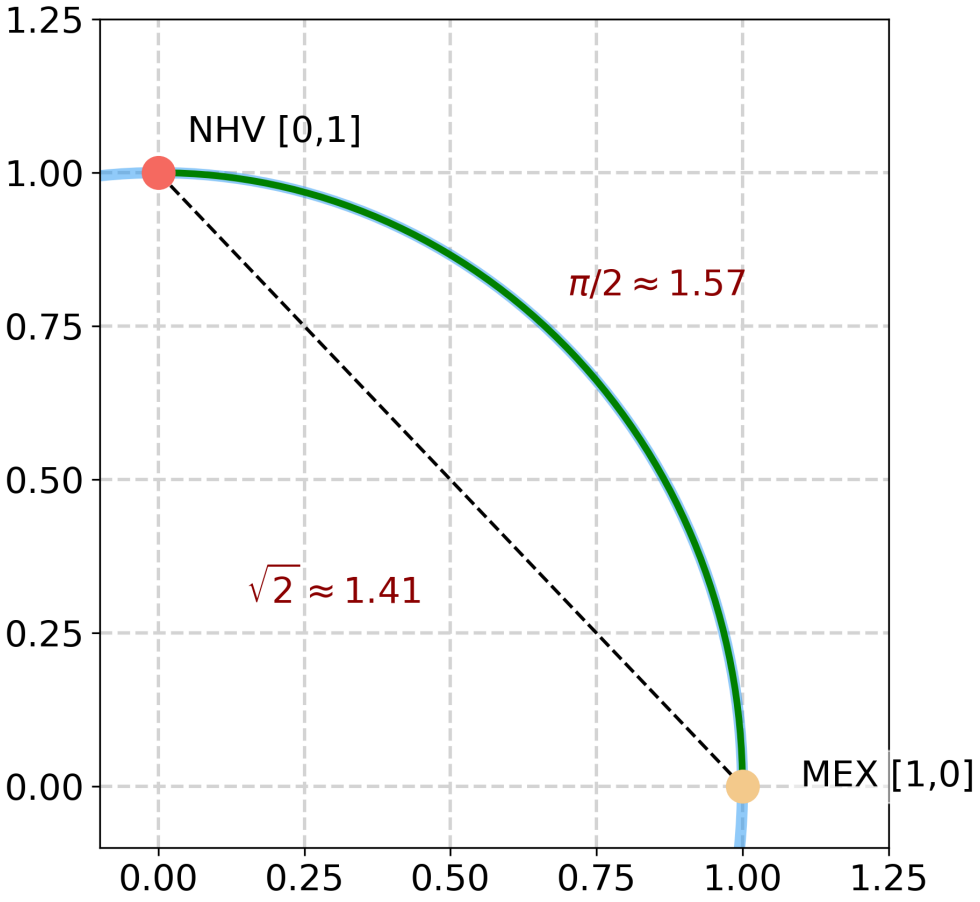

Is a circle locally Euclidean?

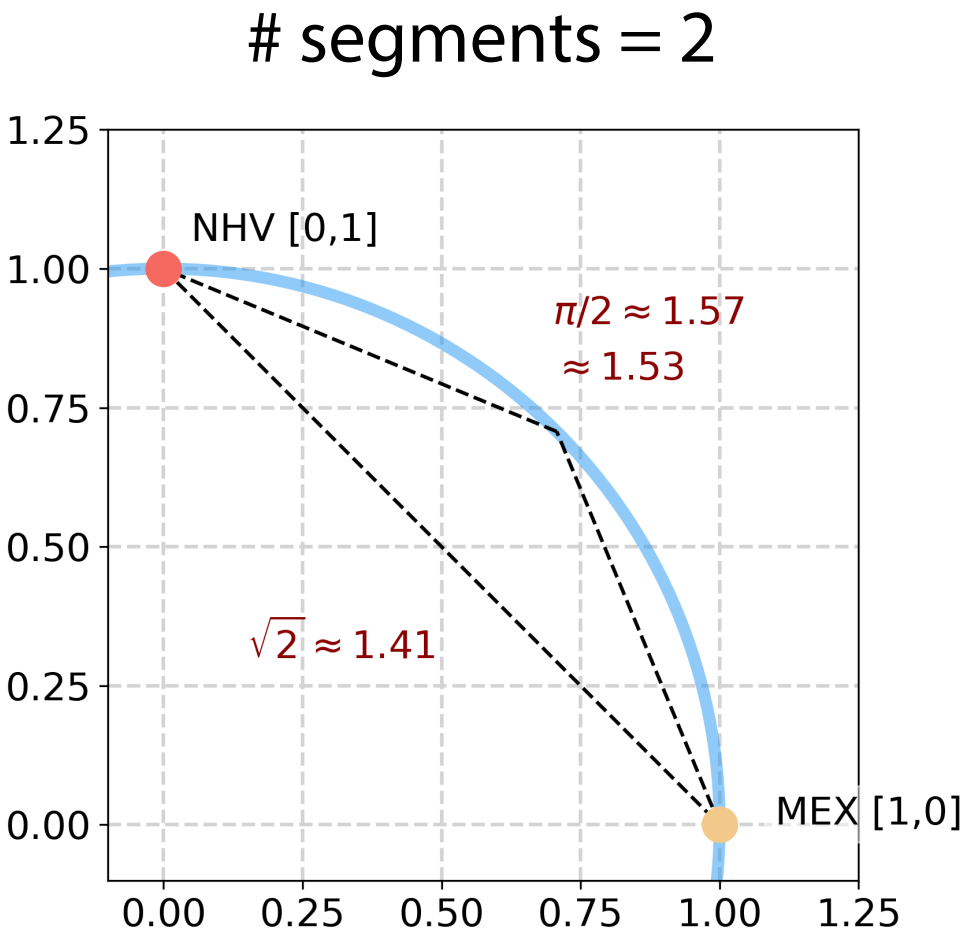

Is a circle locally Euclidean?

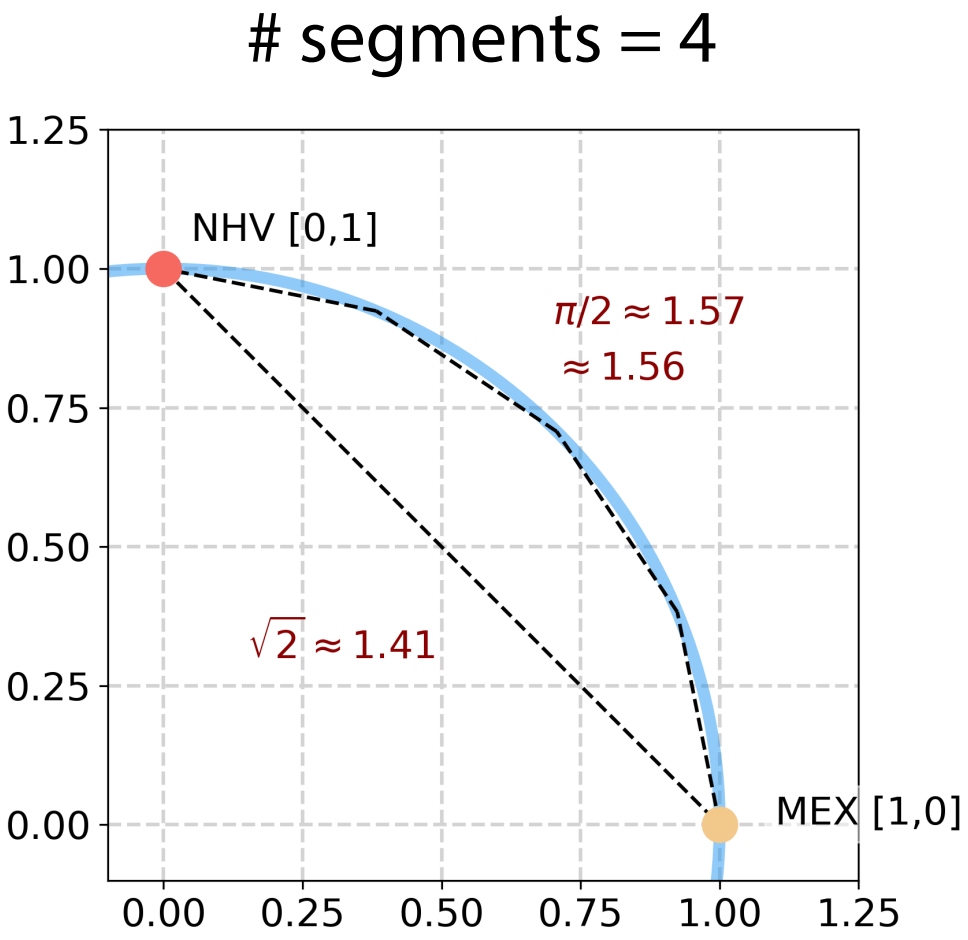

Is a circle locally Euclidean?

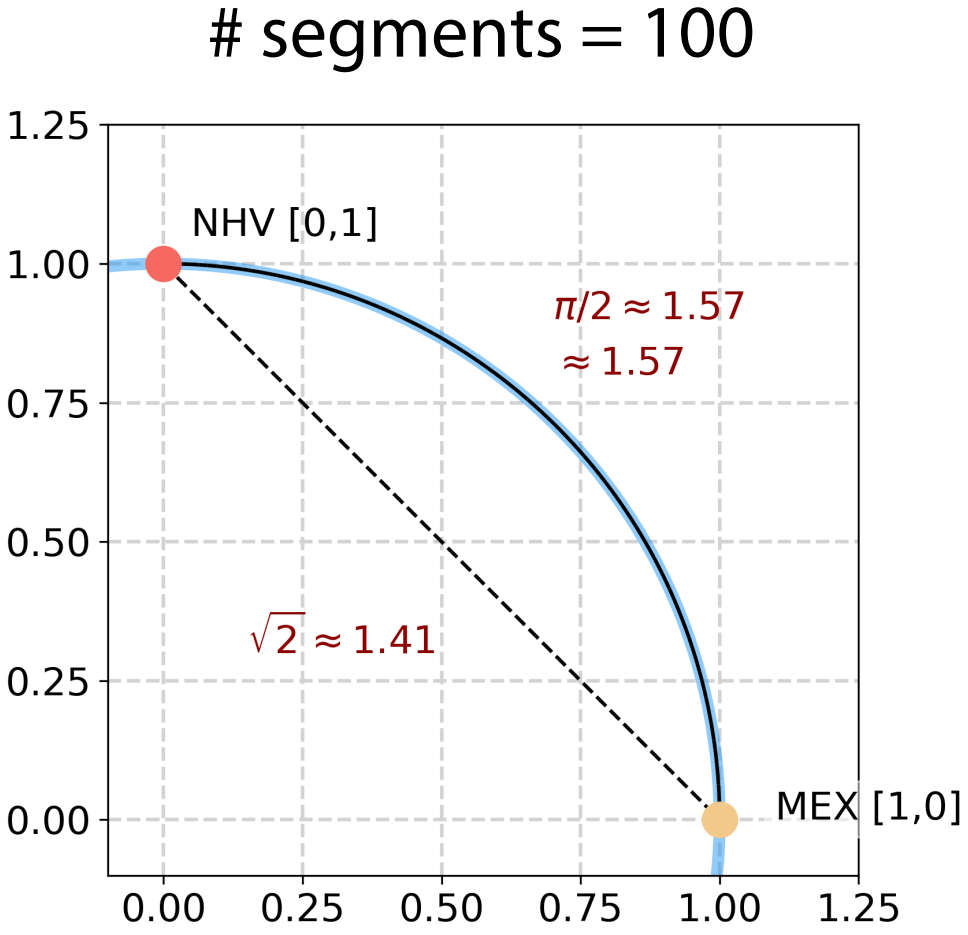

Is a circle locally Euclidean?

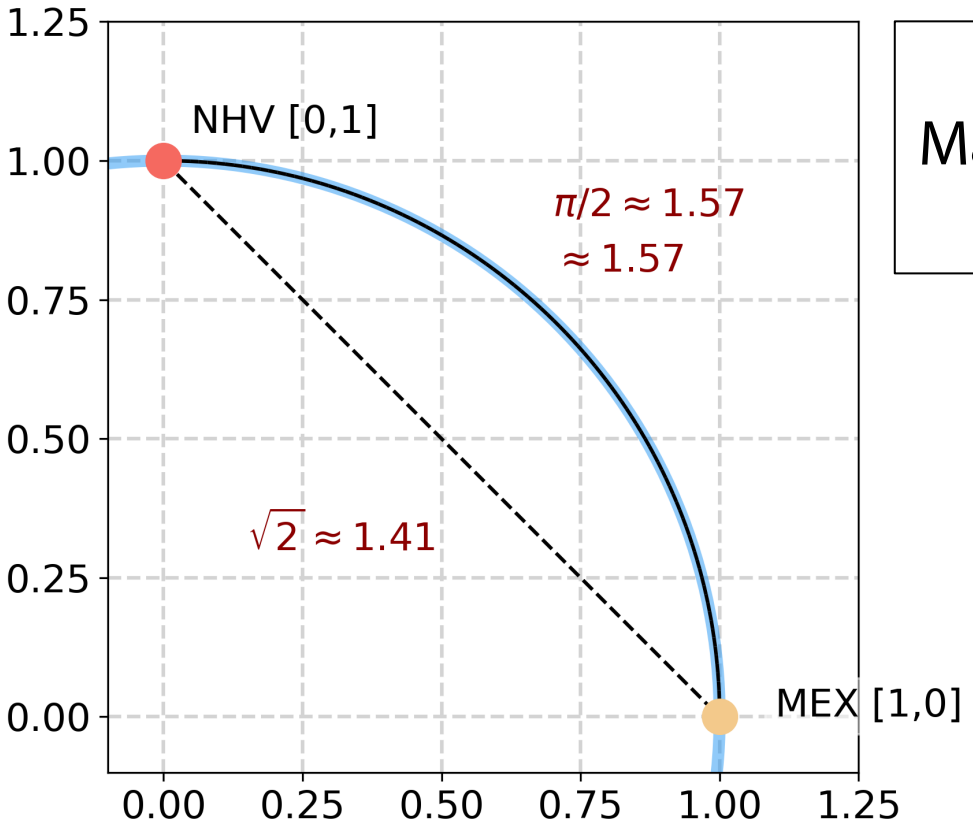

Key Insight:
Manifolds are **locally** Euclidean,
but not **globally** Euclidean

Is a circle locally Euclidean?

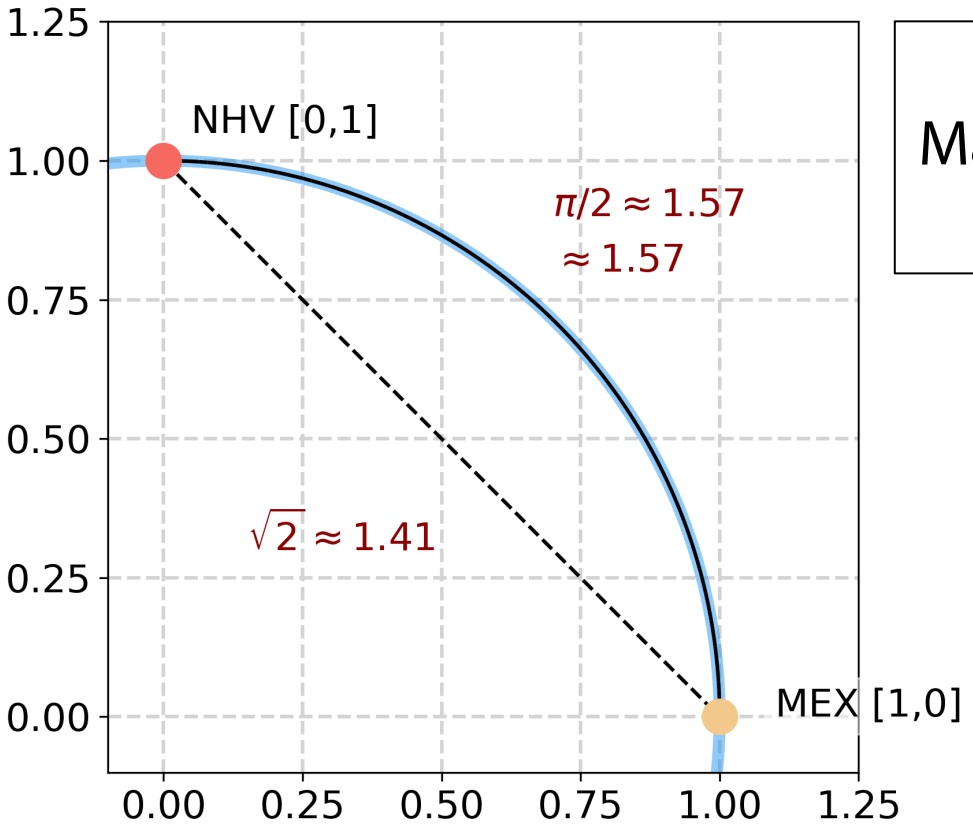

Key Insight:
Manifolds are **locally** Euclidean, but not **globally** Euclidean

How do we define local versus global when we only have a collection of points?

The swiss roll is a plane ($\mathbb{R}^2$) embedded in $\mathbb{R}^3$

True manifold

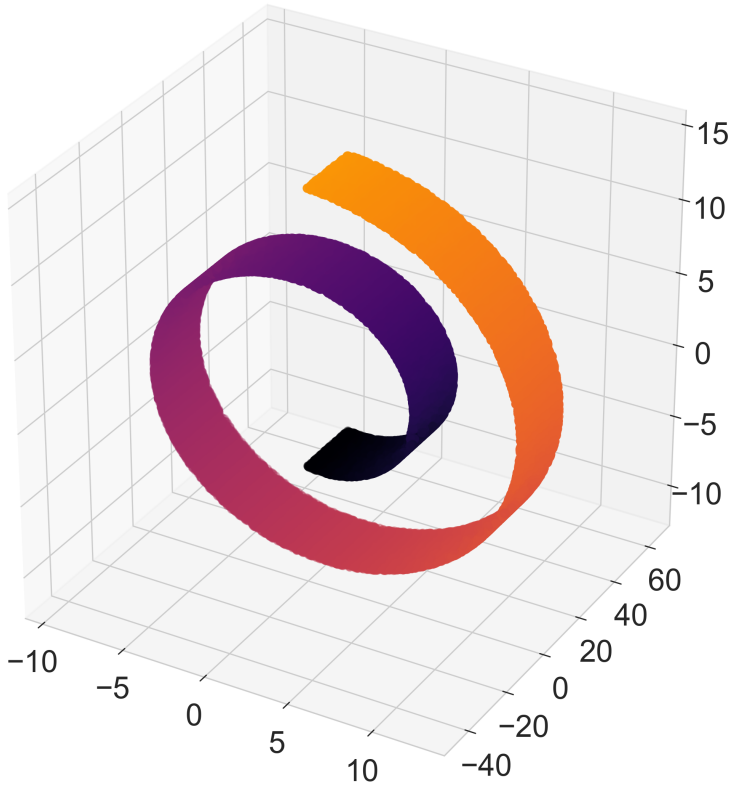

Data sampled from manifold

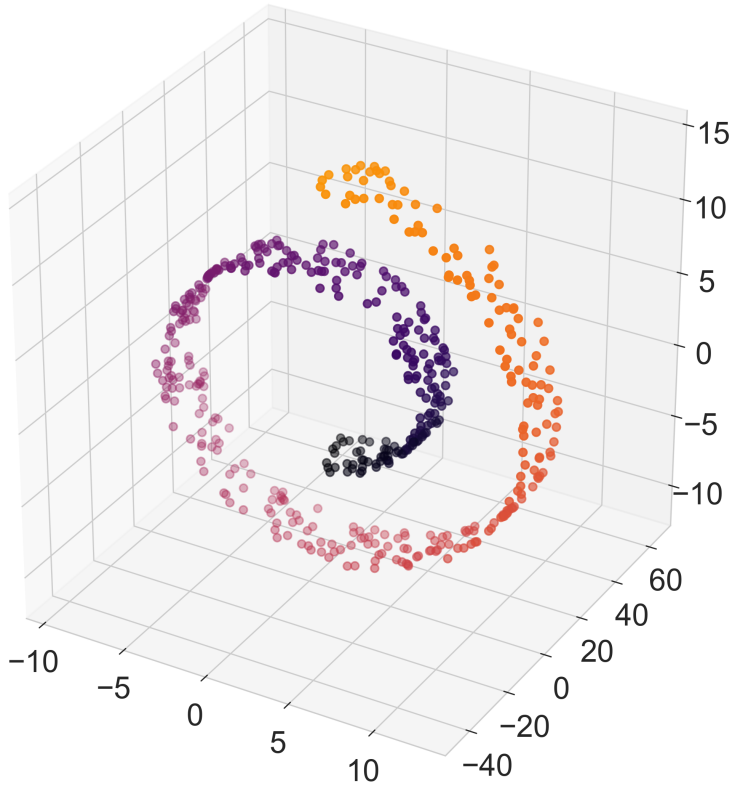

How do we quantify close (local) versus **far** (**global**) in a dataset?

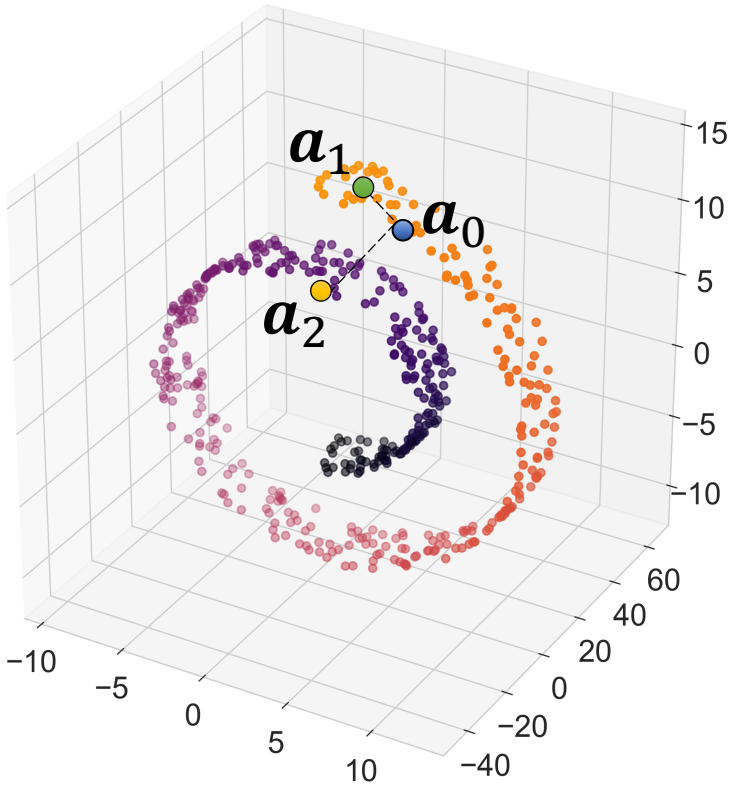

How do we quantify close (local) versus **far** (global) in a dataset?

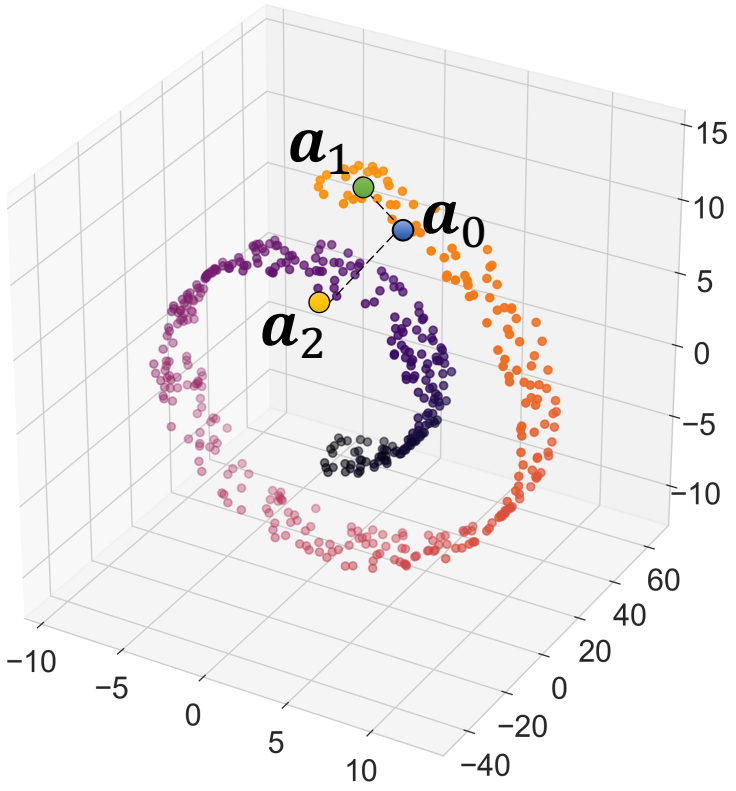

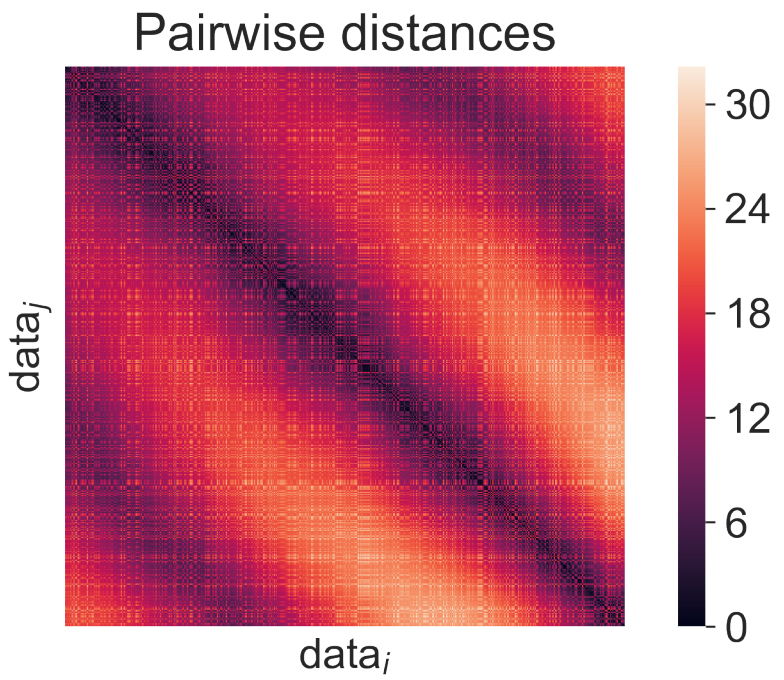

How do we quantify close (local) versus **far** (global) in a dataset?

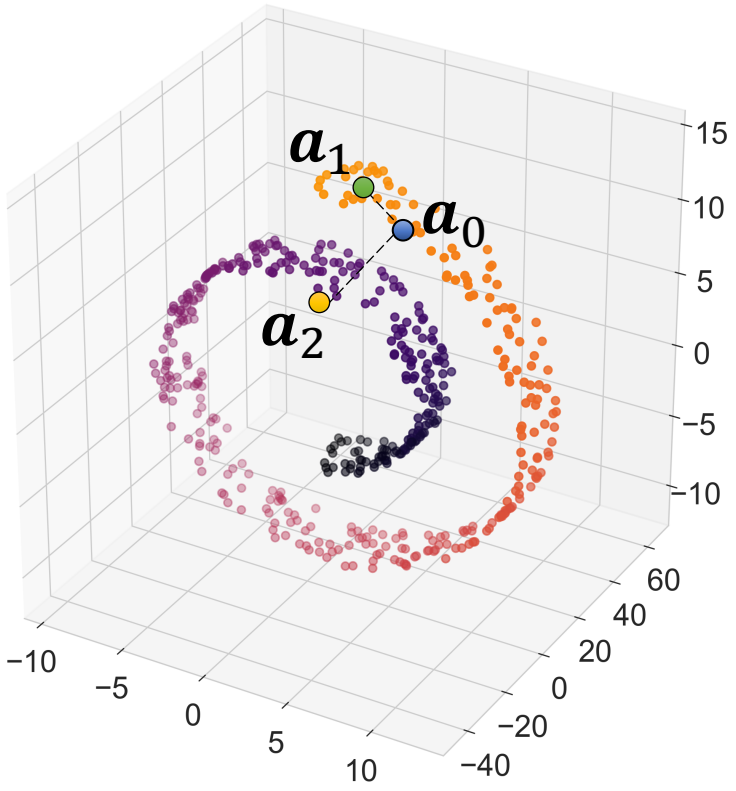

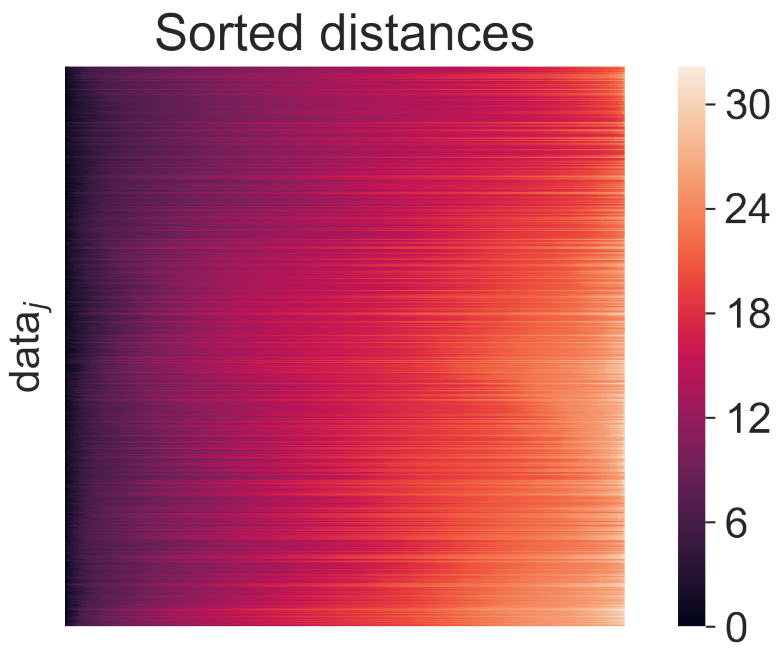

How do we quantify close (local) versus far (global) in a dataset?

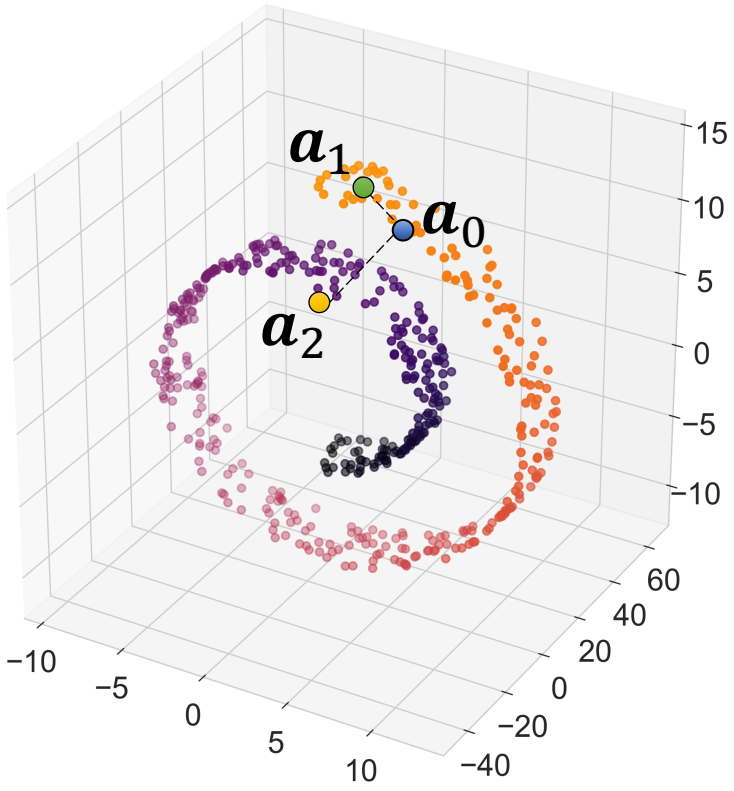

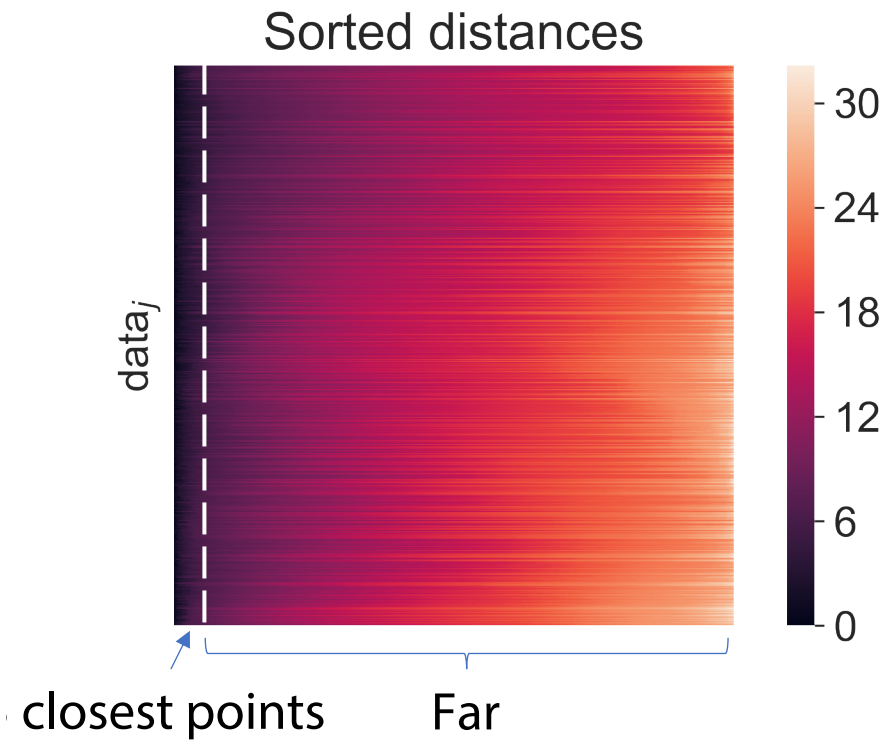

How do we quantify close (local) versus **far** (**global**) in a dataset?

Point out that A2 would be “sort of close” in Euclidean space

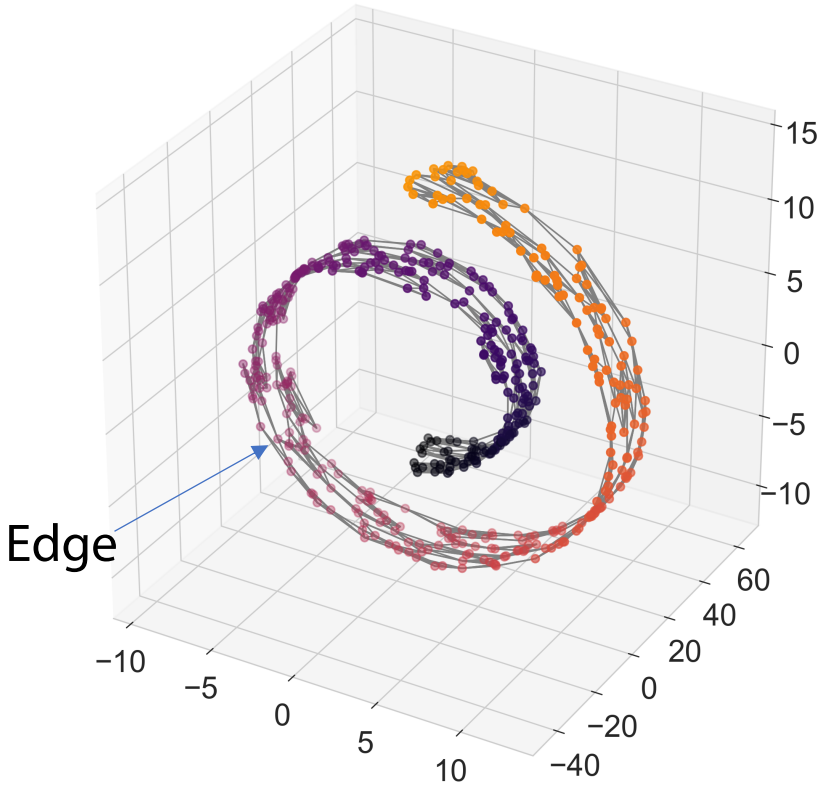

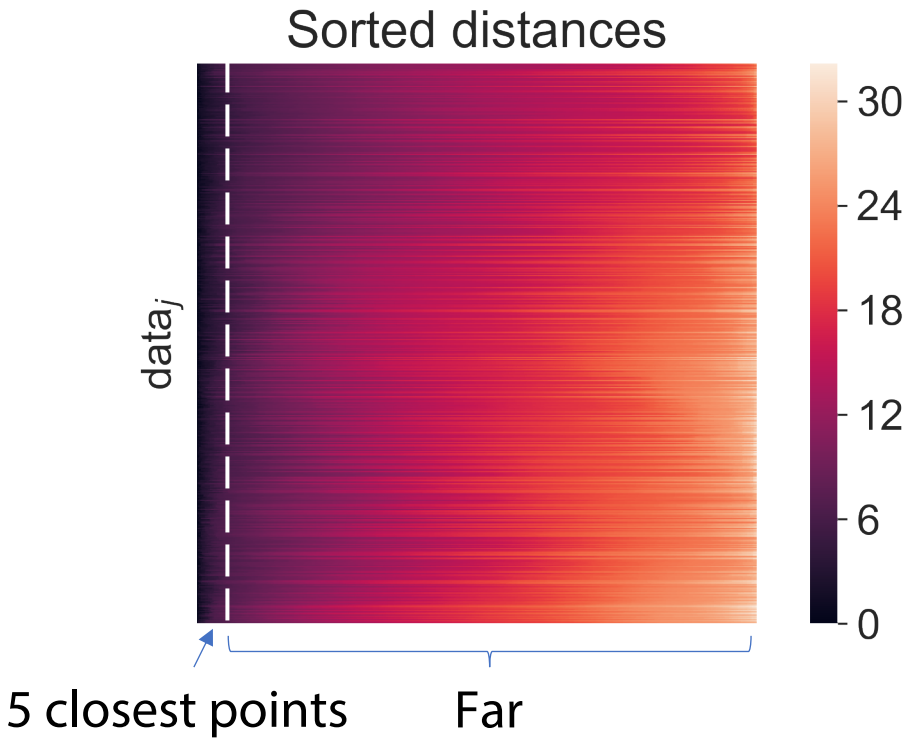

What is a “graph”?

Plot

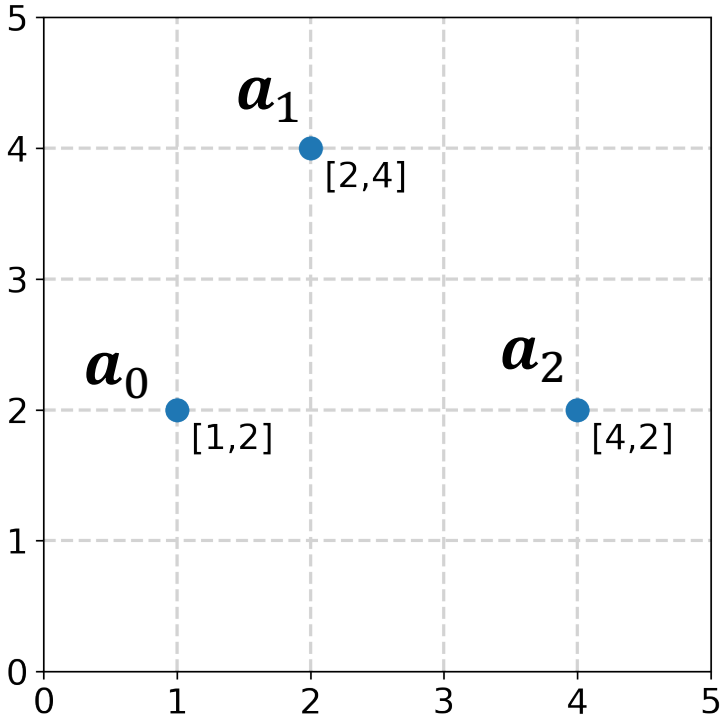

Graph

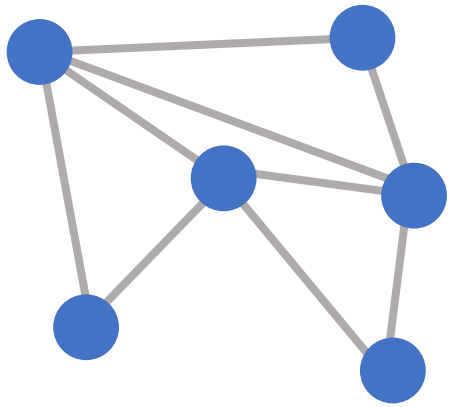

$G = (V, E)$

● = V = vertices / nodes

— = E = edges

What is a “graph”?

Things that can be modelled as a graph

Thing	Vertices	Edges
Internet	Computers	Network connections
Traffic	Intersections	Roads
Social network	People	Friendships
Cell similarities	Cells	Similarity relationships

Graph

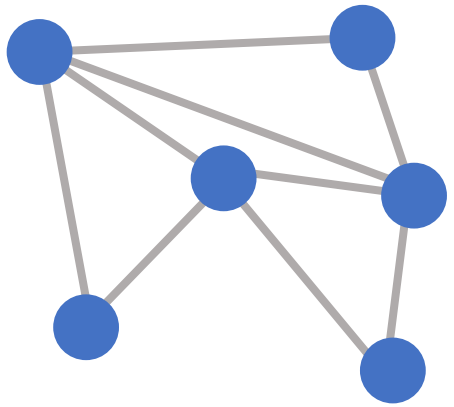

$G = (V, E)$

● $= V =$ vertices

— $= E =$ edges

Graph walks approximate manifold distances

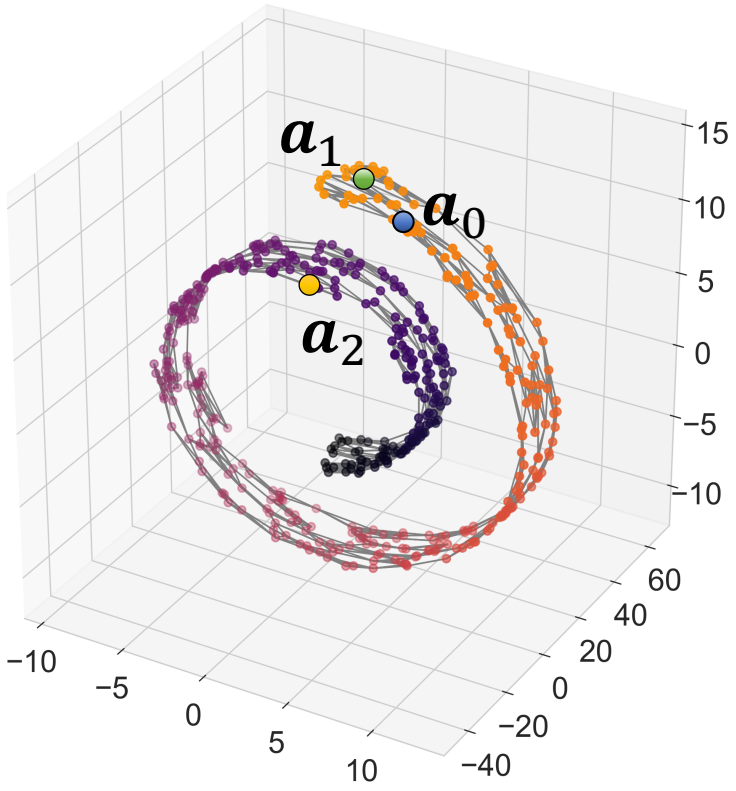

Shortest path between points

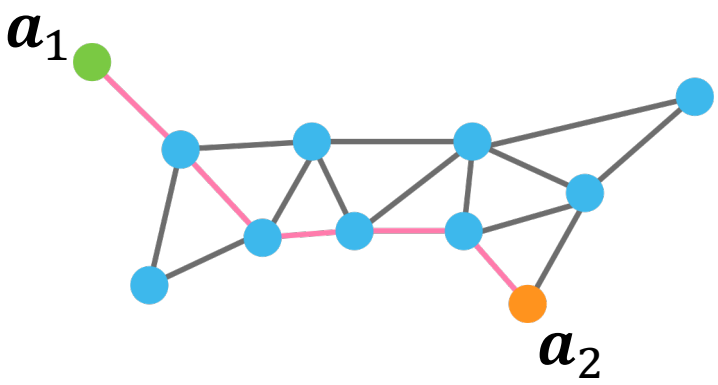

Graph walks approximate manifold distances

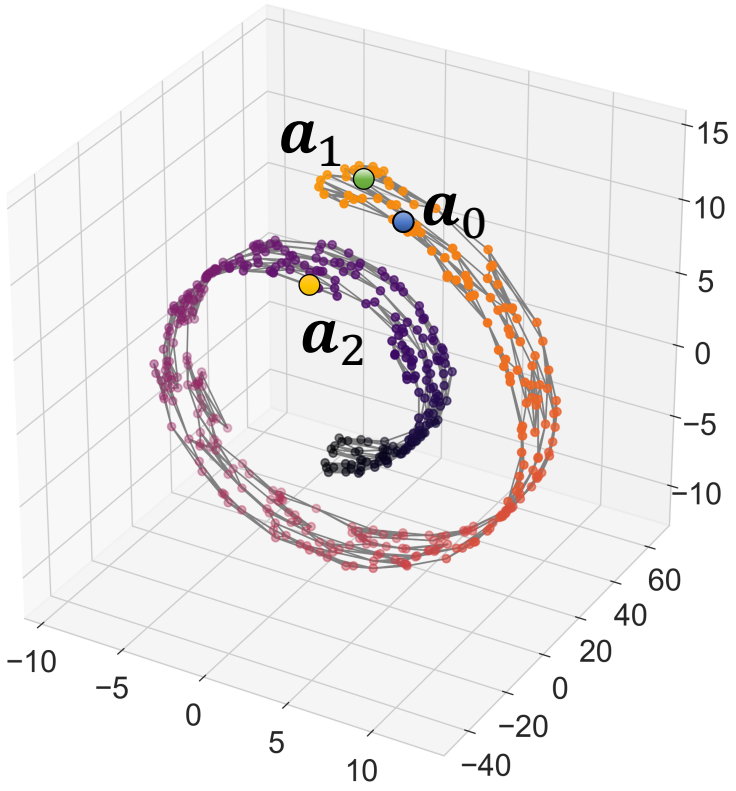

Shortest path between points

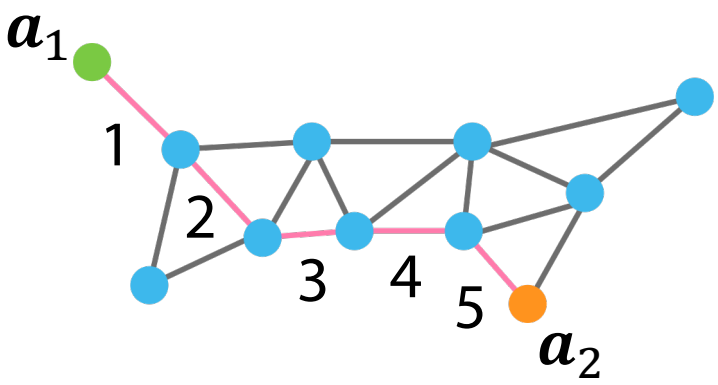

$$d_{geodesic}(a_1, a_2) = \text{shortestpath}(a_1, a_2)$$
$$= 5$$

Graph walks approximate manifold distances

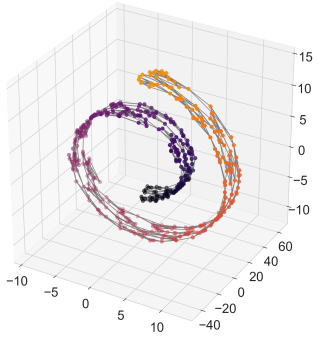

k-NN Graph Geodesic Distances

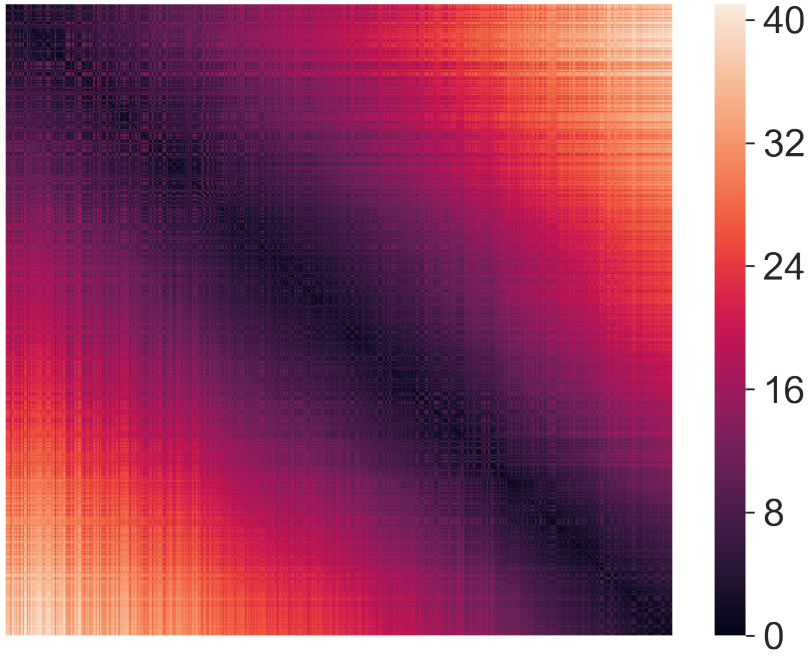

Manifold Distances

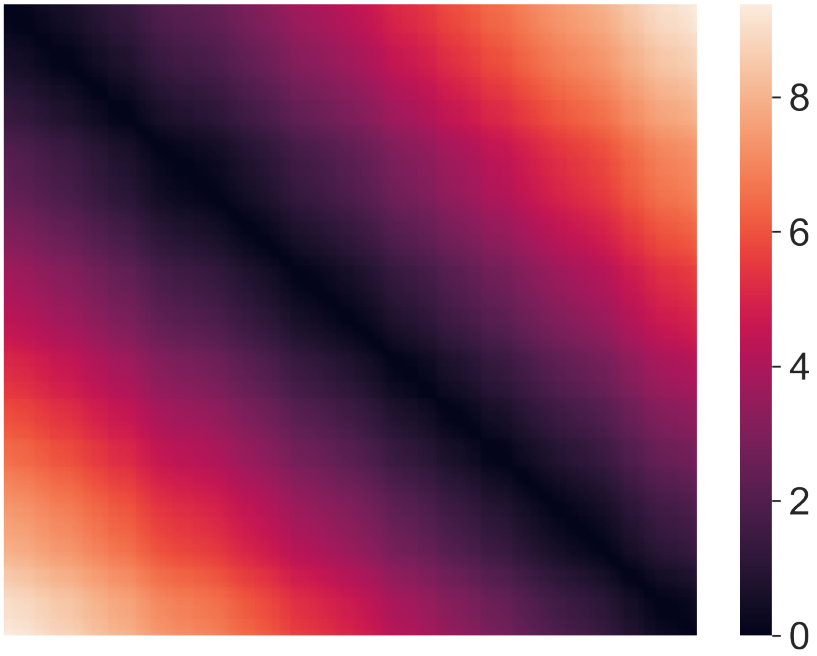

What if the data isn’t entirely connected?

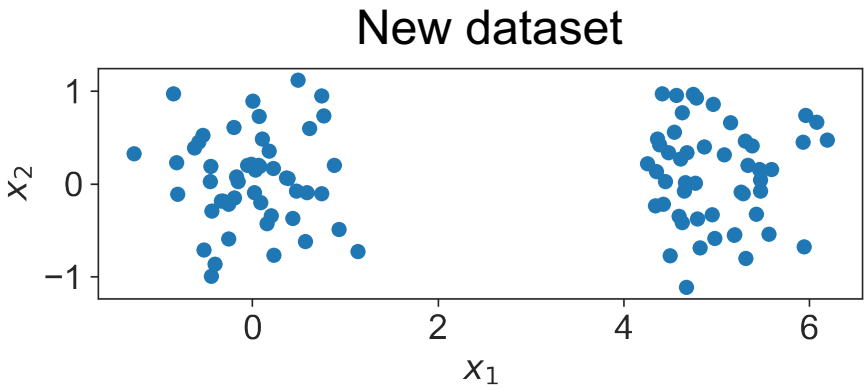

What if the data isn't entirely connected?

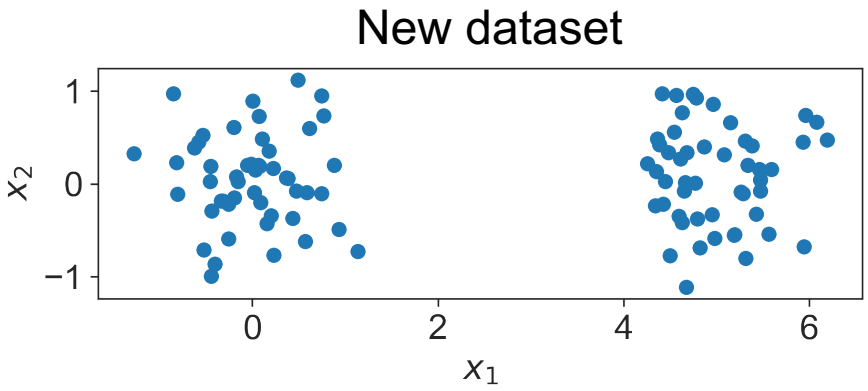

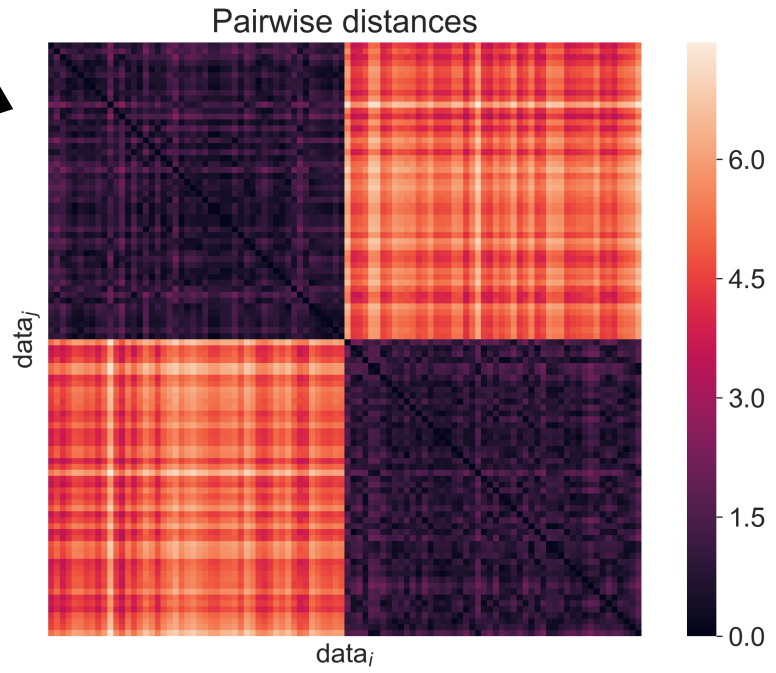

What if the data isn't entirely connected?

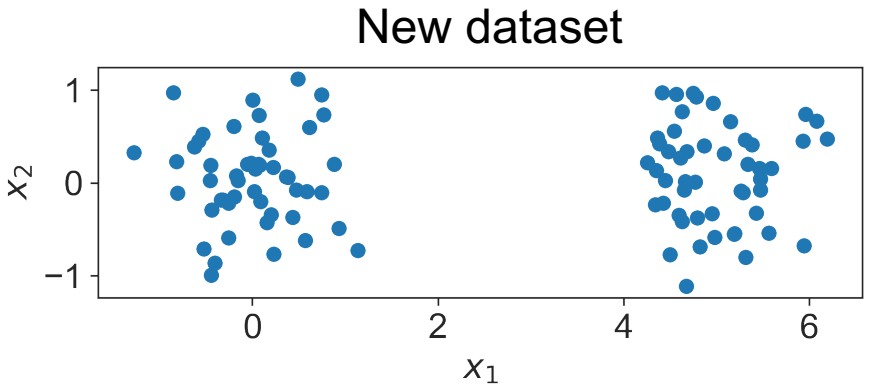

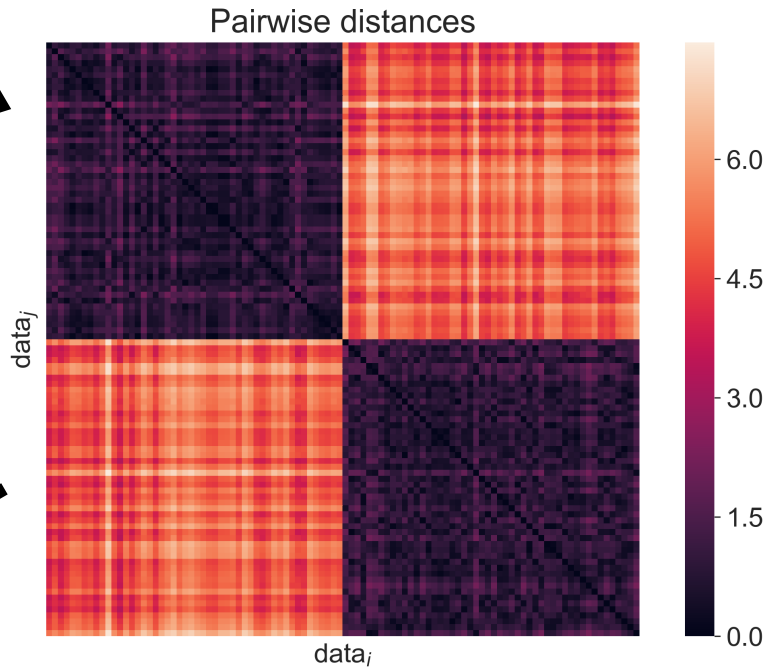

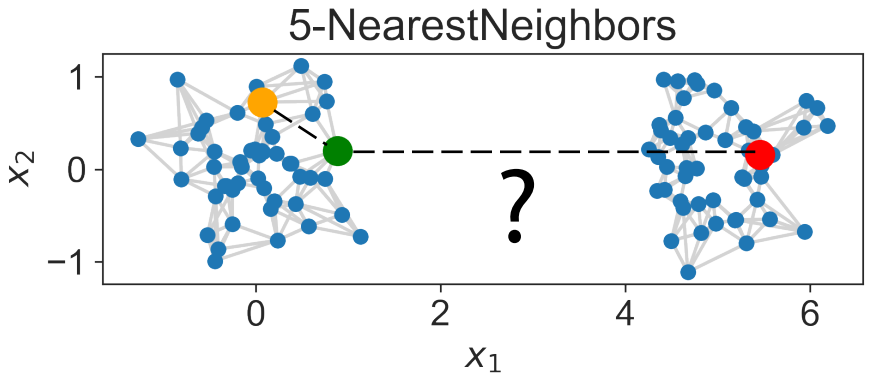

What if the data isn't entirely connected?

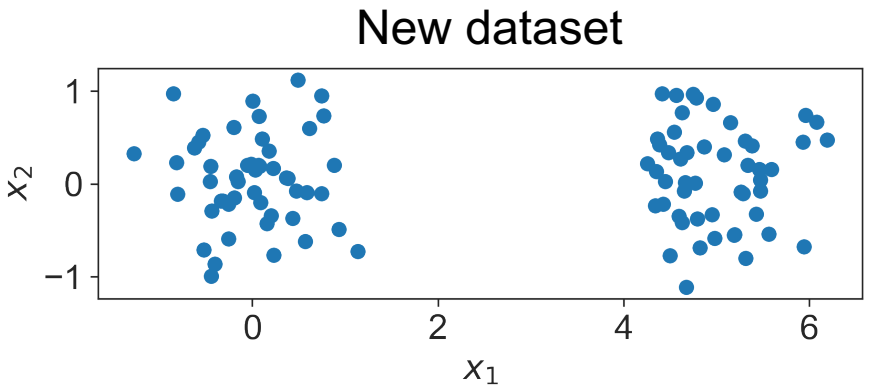

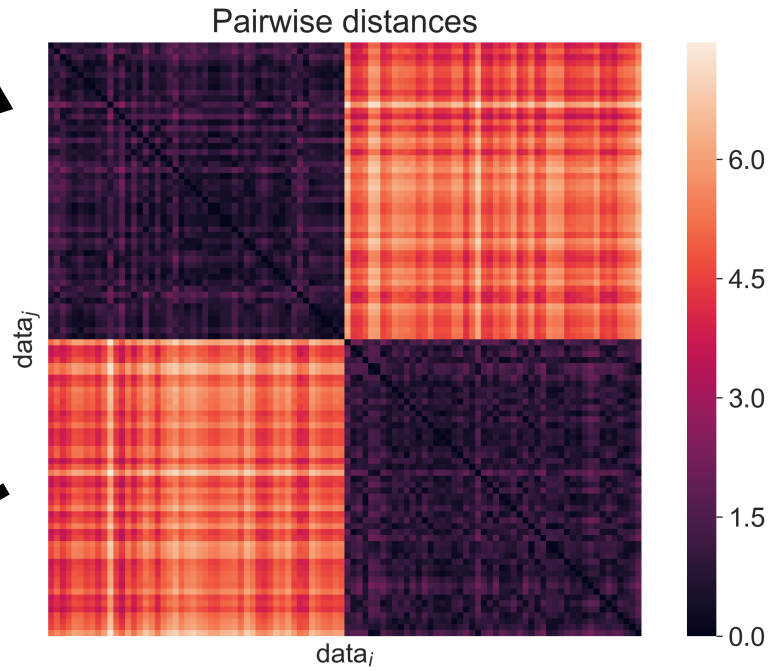

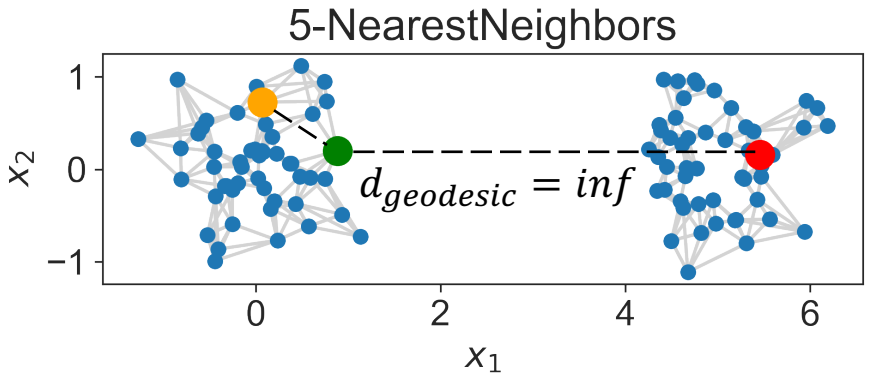

What if the data isn't entirely connected?

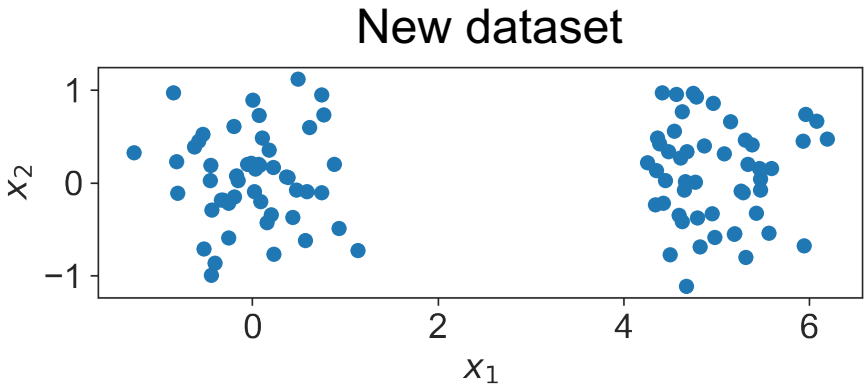

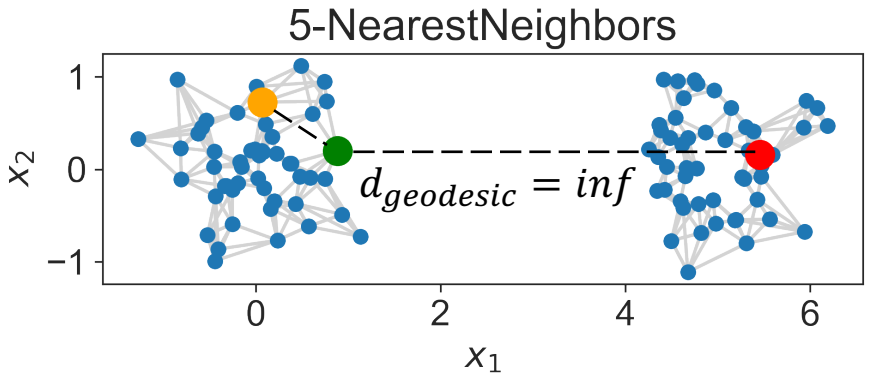

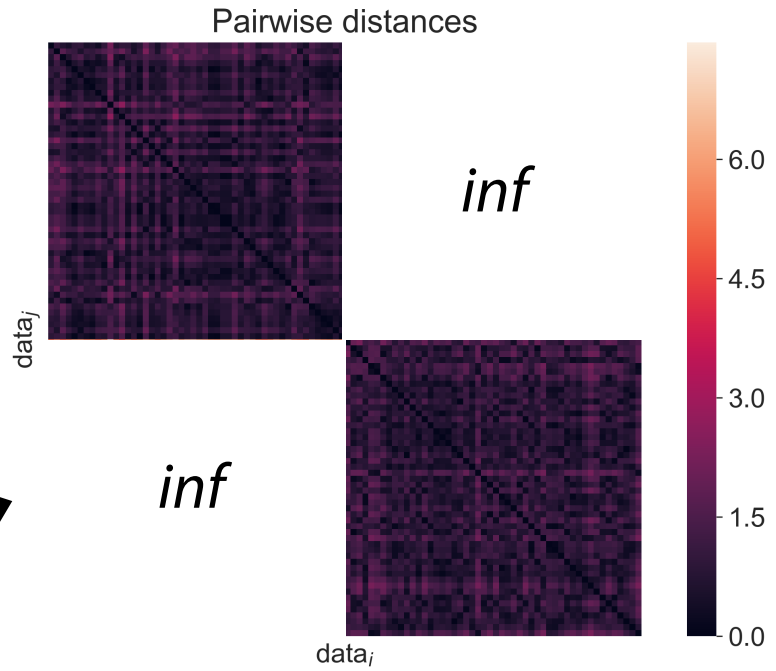

Summary: learning graphs from data

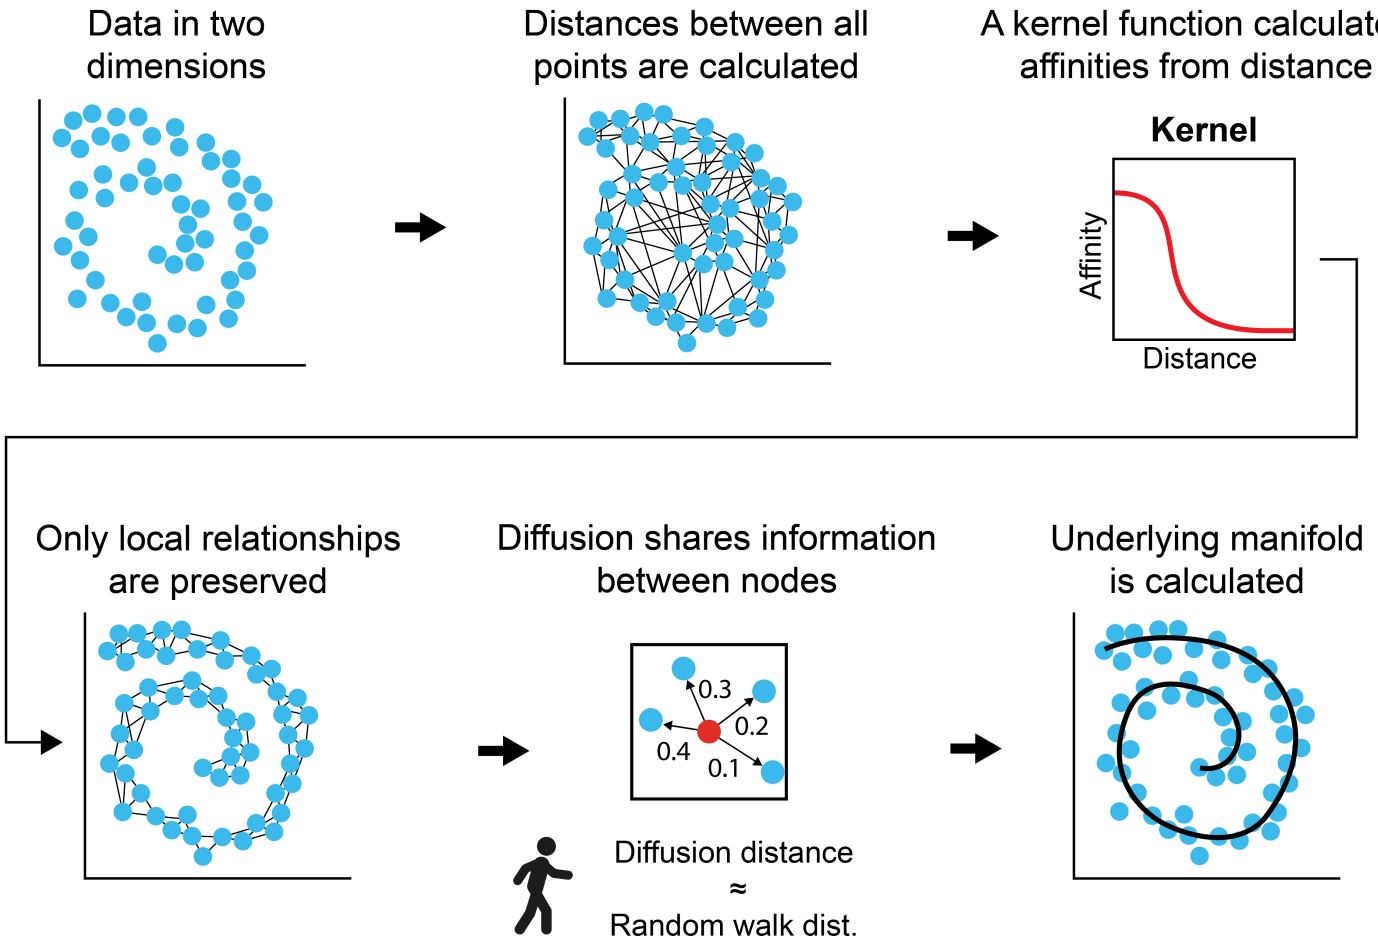

Summary: learning graphs from data

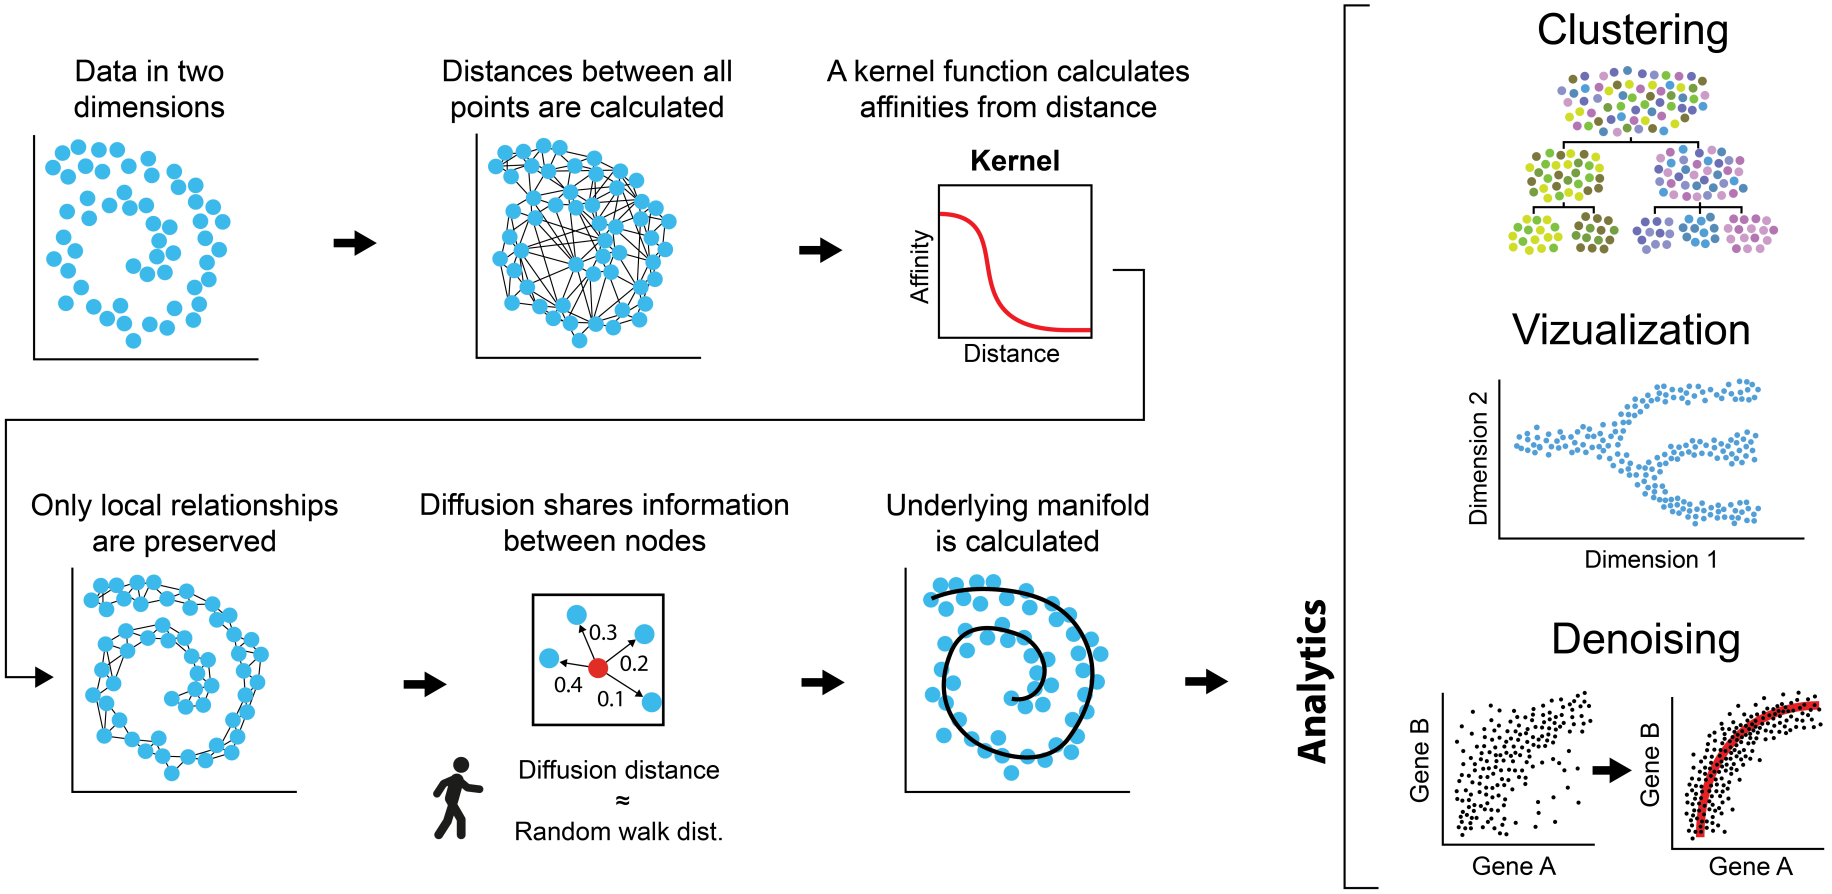

Matrix representations of graphs

G

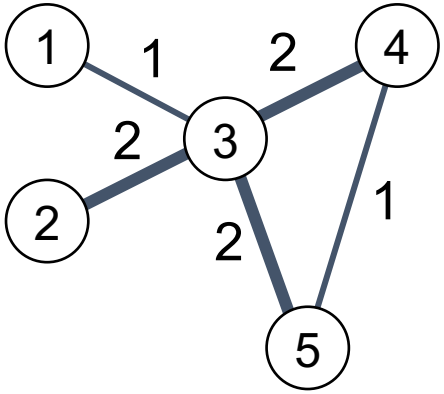

A adjacency matrix

$$\mathbf{A} = \begin{vmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 \\ 1 & 2 & 0 & 2 & 2 \\ 0 & 0 & 2 & 0 & 1 \\ 0 & 0 & 2 & 1 & 0 \end{vmatrix}$$

D degree matrix

$$\mathbf{D} = \begin{vmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 7 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 3 \end{vmatrix}$$

Laplacian matrix

$$\mathcal{L} = \mathbf{D} - \mathbf{A} = \begin{vmatrix} 1 & 0 & -1 & 0 & 0 \\ 0 & 2 & -2 & 0 & 0 \\ -1 & -2 & 7 & -2 & -2 \\ 0 & 0 & -2 & 3 & -1 \\ 0 & 0 & -2 & -1 & 3 \end{vmatrix}$$

→

- Creating visualizations
- Clustering
- Denoising data associated with the nodes

Exercise – choosing k

k = 5

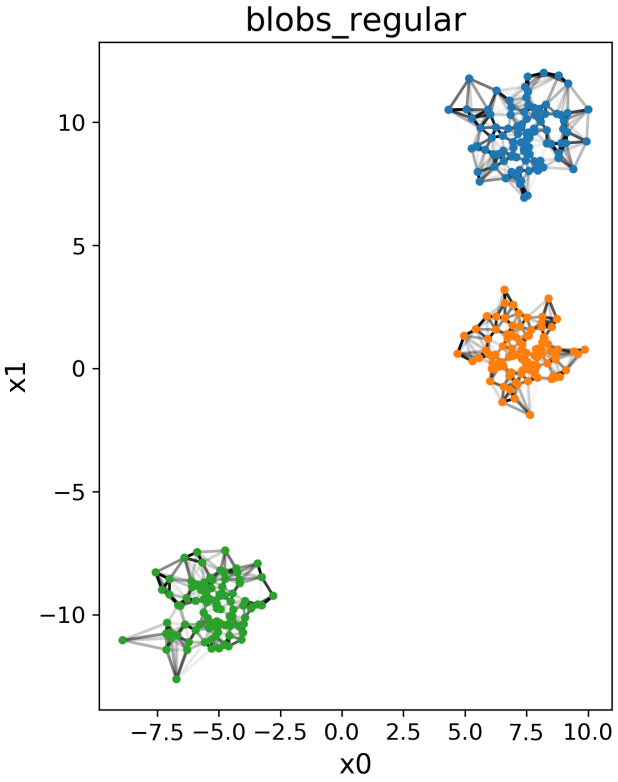

k = 10